

# Managing Large Databases



## Using MongoDB

---

# **Table of Contents**

---

## **Sessions**

- Session 1 – Introduction to MongoDB**
- Session 2 – MongoDB Databases**
- Session 3 – MongoDB Operators**
- Session 4 – Aggregation Pipeline**
- Session 5 – Database Commands**
- Session 6 – MongoDB Shell Methods**
- Session 7 – MongoDB Indexing**
- Session 8 – MongoDB Replication and Sharding**
- Session 9 – Transaction Management**
- Session 10 – MongoDB Tools**
- Session 11 – MongoDB Cloud**
- Session 12 – MongoDB Database Connectivity With Python**
- Appendix – Case Studies**



## SESSION 1

### INTRODUCTION TO MONGODB

---

#### Learning Objectives

In this session, students will learn to:

- Explain various types of databases
- Describe how data is stored in the MongoDB Server
- Explain the installation of MongoDB Server and MongoDB Shell

Databases have been in use for quite some time now. Most of the companies use structured databases, such as SQL Server and Oracle, to store information related to their products, customers, sales, employees, and so on. These databases grow continuously with more and more new information. Companies must invest in the required hardware to store the ever-growing data. As the databases explode, some of the existing database systems become very expensive to upgrade and maintain. In addition, the performance of these databases is affected due to large volume of data.

To overcome the issues that structured databases posed with increasing data, companies are now opting for NoSQL databases. Especially companies that host social media apps or gaming apps, which encounter an ever-increasing customer or player base. MongoDB is one of the popular NoSQL databases.

This session will provide an overview of the types of databases and the prevalence of NoSQL databases. This session also explains the architecture of MongoDB and illustrates the steps to install the MongoDB server and shell.

## 1.1 Databases

---

Data refers to any information and a database is a collection of data. An example of a database is a telephone directory or a product catalog. Depending on the way the data is stored, there are different types of databases. These are:

<b>Hierarchical</b>	Similar to a family tree, this type of database stores data in the form of parent-child nodes. <b>Example:</b> Windows Registry
<b>Relational</b>	This type of database stores data in the form of rows and columns that together form a table (relation). <b>Example:</b> MySQL, Oracle
<b>Object-oriented</b>	This type of database stores data in the form of objects similar to the objects used in an object-oriented programming language. <b>Example:</b> PostgreSQL
<b>Network</b>	This database stores data as multiple children and parent nodes that results in complex database structures. <b>Example:</b> RDM Server
<b>NoSQL</b>	This database does not use a relational model and is used for analyzing large sets of distributed data. <b>Example:</b> MongoDB

Before getting into the details of MongoDB, let us understand the difference between relational databases, which are quite popular, and NoSQL databases, which are becoming more and more popular.

Table 1.1 distinguishes the features of relational and NoSQL databases.

<b>Attribute</b>	<b>Relational Database</b>	<b>NoSQL Database</b>
Data storage	It stores data within rows and columns in tables.	It stores data in various formats, such as documents, graphs, or key-value pairs.
Schema	It requires every table in the database to have a defined schema and the schema cannot be altered after it is created.	It does not require data to have a defined schema. It supports dynamic schema, which allows alterations to the storage structure as required.
Scaling	To support increase in data volume and database users, the hardware of the server must be updated with RAM or processors, or a completely new server machine must be installed. Adding more resources to a server is referred to as vertical scaling.	To support increase in data volume and database users, multiple computers or servers with minimal configuration can be installed. Adding more servers that contain same set of databases with different sets of information is referred to as horizontal scaling.
Performance	The information related to a single entity is stored in multiple tables, based on normalization. Therefore, if data is to be retrieved for a single entity, multiple tables must be iterated through. This results in a decrease in performance, especially if the volume of data is large.	The information related to a single entity is stored together in a single location. Therefore, if data is to be retrieved for a single entity, a simple query can be used. This improves the performance of the database.
Consistency	Relational DataBase Management System (RDBMS) follows the Atomicity, Consistency, Isolation, and Durability (ACID) properties.	NoSQL Database does not follow the ACID properties.

Attribute	Relational Database	NoSQL Database
Application	RDBMS is a best fit when data items must be related and stored in a rules-based, consistent manner. Examples include e-commerce transactions and stock tracking.	NoSQL database is a best fit for applications that require huge data volumes and low latency. Examples include online gaming, social media, and shopping apps.

**Table 1.1: Relational Databases vs. NoSQL Databases**

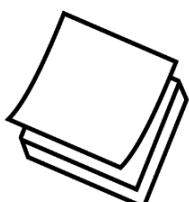
## 1.2 Overview of MongoDB

---

Consider a social media mobile app where a movie star with a million followers has posted his latest movie stills. Several thousands of his followers are commenting on the movie stills and the movie star is also responding to these comments. Comments and responses are a mix of text messages, animations, pictures, and emojis. If this social media app is placed on RDBMS, there would be one table each to store:

- the details of the movie star,
- the details of the followers,
- the posts put up by the movie star, and
- the comments and responses.

All these tables will be related to each other based on one or more columns. In this case, the time taken for reading/writing such large, changing, and varied data to the server will be really long. Consequently, the user experience of the mobile app would not be very good. However, a NoSQL database would perfectly fit in such a scenario because all the data can be stored together. This data storage will speed up data retrieval leading to good user experience. One example of a NoSQL database is MongoDB.



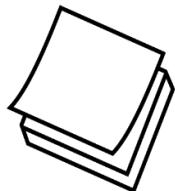

---

Depending on how data is stored in the database, there are four types of NoSQL databases:

- Document databases: Store data as JSON, BSON, or XML documents
- Key-value databases: Store data with a label or attribute and the value
- Wide-column databases: Store data in columns

- Graph databases: Store each data element a node and relationship between the data elements as links or relationships
- 

MongoDB is a NoSQL database program that stores information in formats similar to JavaScript Object Notation (JSON) called the Binary JavaScript Object Notation (BSON). MongoDB is an open-source software which works with large sets of distributed data. It provides high performance, high availability, and automatic scaling.



---

The JSON format is a standard human-readable format used to transmit data in the form of key-value pairs. BSON is used to store and access documents transmitted using JSON format. MongoDB stores documents in BSON format.

---

### 1.2.1 MongoDB Architecture

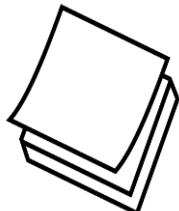
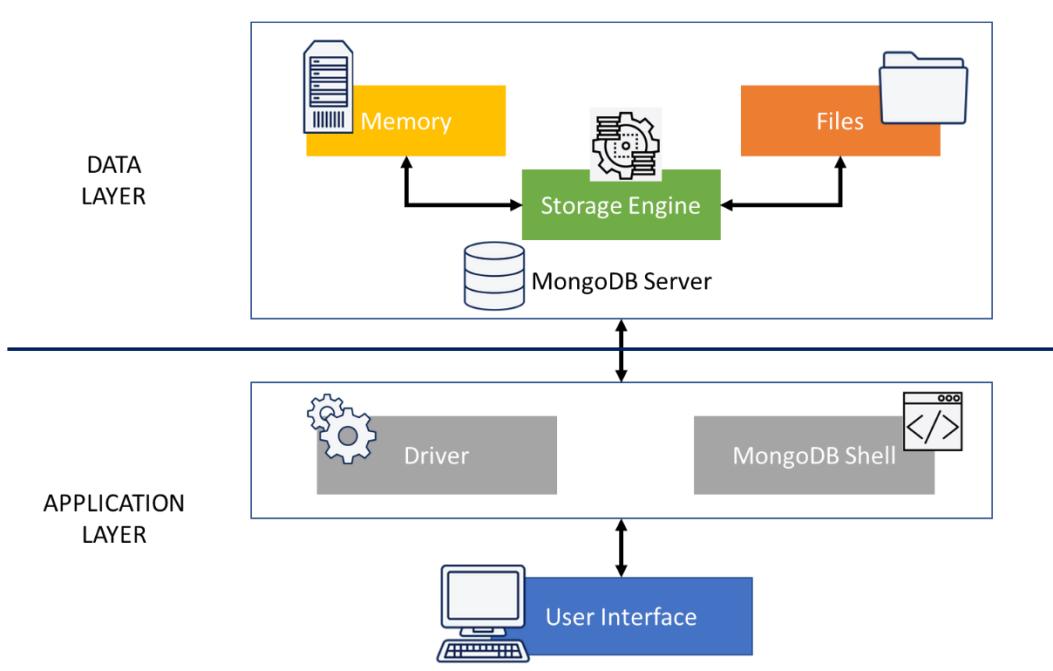
MongoDB operates in two layers: the application layer and the data layer.

The application layer is made up of:

- User interface that the user uses to interact with the database by using a mobile app (Android or iOS) or directly from the Web,
- MongoDB driver that helps to make the connection between the user interface and the MongoDB server, and
- MongoDB Shell which is the command-line interface that users can use to interact with the MongoDB server.

The data layer consists of the MongoDB server that receives the queries from the application layer and sends them on to the storage engine. The storage engine then reads or writes the data to the memory or the file. The storage engine determines how data is stored and the amount of data stored on the disk (file) and memory.

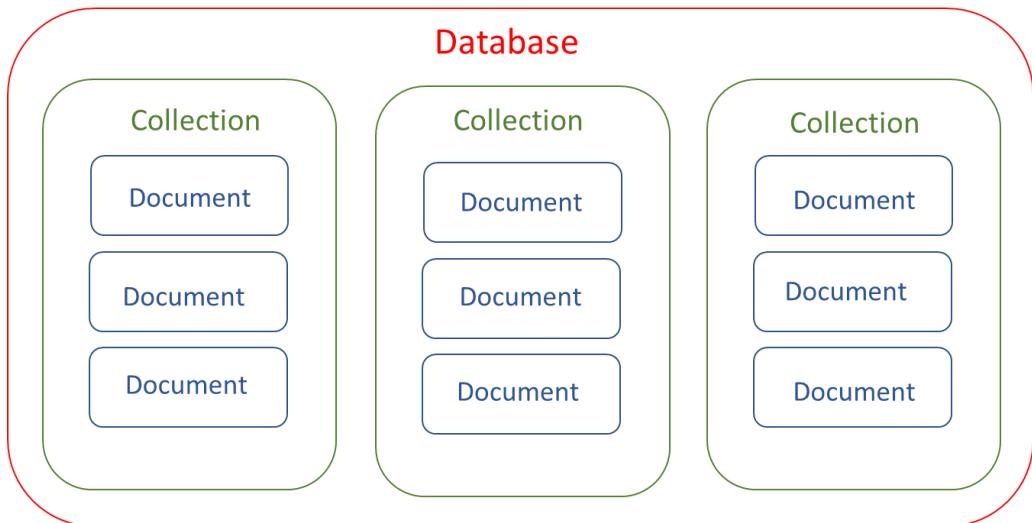
Thus, the vital components of MongoDB that are necessary for its function are:



- MongoDB supports drivers for C, C++, .Net, Go, C#, Python, Node.js, and so on.
- WiredTiger is the default MongoDB storage engine. However, MongoDB also uses other storage engines, such as the in-memory storage engine, MMAPv1 storage engine, and the encrypted storage engine.

### 1.2.2 Database, Collection, and Document

Data is stored in documents (rows and columns in RDBMS) on the MongoDB Server. Multiple documents together make up a collection (tables in RDBMS). Multiple collections make up a database.



MongoDB stores data in the document as BSON documents. BSON is a binary representation of JSON documents, although BSON supports more data types than JSON. The BSON format will look as follows:

```
{
    field1: value1,
    field2: value2,
    field3: value3,
    ...
    fieldN: valueN
}
```

A sample document with values in the fields will look as follows:

```
{
    _id: ObjectId("5099803df3f4948bd2f98391"),
    name: { first: "Larissa", last: "Smith" },
    birth: new Date('Jun 23, 2015'),
    grade: 6,
    marks: [{sub: "English",
              score: 89},
             {sub: "Math",
              score: 75},
             {sub: "Science",
              score: 92},
             {sub: "Social Studies"
              score: 70}],
}
```

The sample document shows key-value pairs which form the basic unit of data in MongoDB.

### 1.2.3 MongoDB Editions

MongoDB is available both as a server and as a service on the cloud. As a server, MongoDB is available as Community Server and Enterprise Server. While Community Server is a free edition, Enterprise Server is a commercial edition that is available on subscription. While the core server features are the same across the two editions, the Enterprise edition offers additional operational and management capabilities, an in-memory storage engine, and advanced security features.

As a cloud service, MongoDB is available as MongoDB Atlas. It offers Data as a Service (DaaS) hosting data in the cloud and provides all the features of cloud computing including high availability and scalability.

MongoDB Atlas includes products such as Atlas Database, Atlas Search, Atlas Data Federation, Atlas Charts, Atlas App Services, Atlas Device Sync, Atlas Data Lake, and so on.

With MongoDB Atlas, users can:

- Save time and cost. MongoDB Atlas includes a fully managed database service that provides access to the database resources as required.
- Use a common API to serve any type of database requests.
- Perform multi-region deployments. Atlas Database is available in more than 80 regions across AWS, Google Cloud, and Azure.
- Focus on building rather than managing. Atlas provides automated infrastructure operations that guarantee availability, scalability, and security compliance.
- Work with data as code. In MongoDB, documents map directly to the objects in the code. As a result, data of any structure can be stored, and the schema can be modified when new features are added to the application.



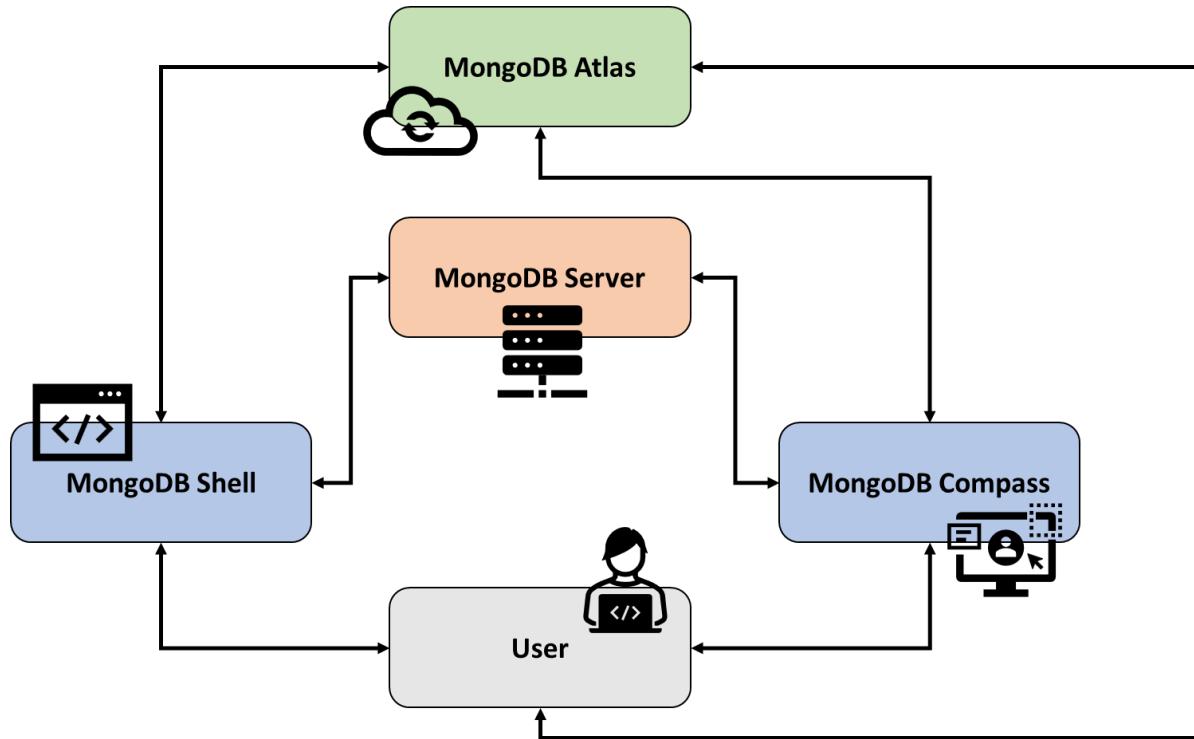
---

MongoDB Atlas allows deployment of databases to any cloud provider rather than providing a single option.

---

The MongoDB Shell is a standalone command-line product that the user can use to connect to the server or the service. MongoDB Compass is Graphical User Interface (GUI) that the users can use to connect to the server or to the service.

Users can also directly access the Atlas platform and work on it without using MongoDB Shell or MongoDB Compass.



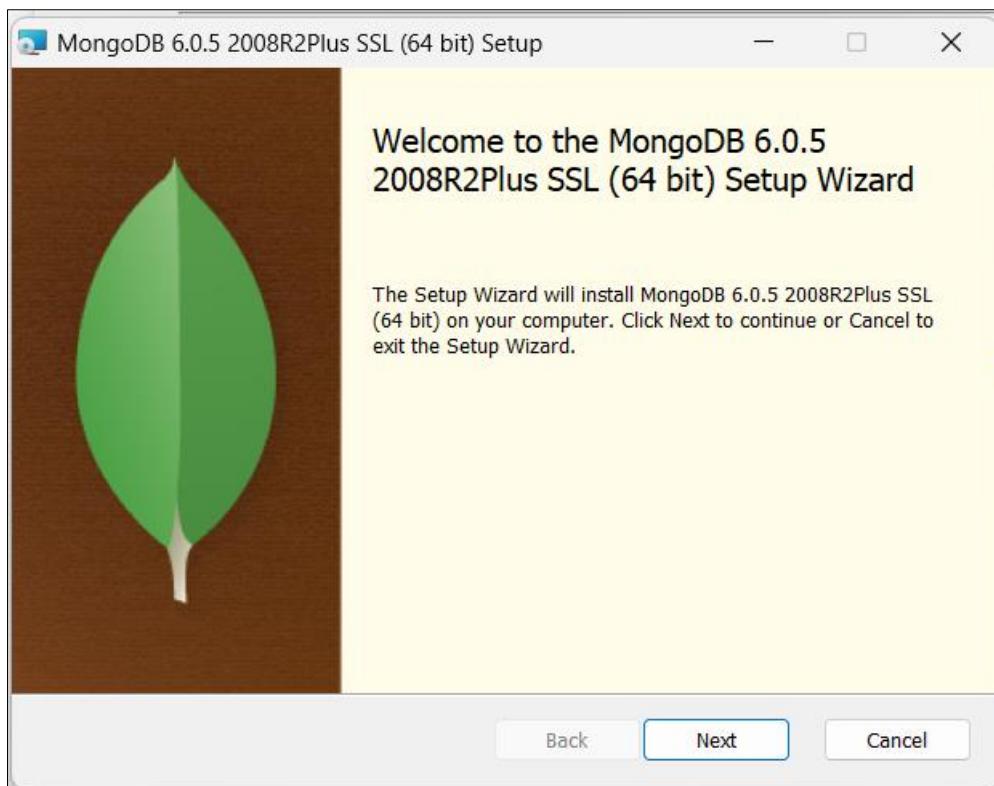
### **1.3 Installing MongoDB Community Server**

This course is based on the MongoDB 6.0 Community edition.

To install the MongoDB 6.0 Community edition on Windows systems:

1. Open a browser and navigate to the URL:  
<https://www.mongodb.com/try/download/community>  
The MongoDB Community Server Download page opens.
2. On this page, scroll down and ensure that:
  - In the **Version** drop-down, **6.0.5(current)** is selected.
  - In the **Platform** drop-down, **Windows** is selected.
  - In the **Package** drop-down, **msi** is selected.
3. Click **Download**.
4. After the download is complete, run the installer.

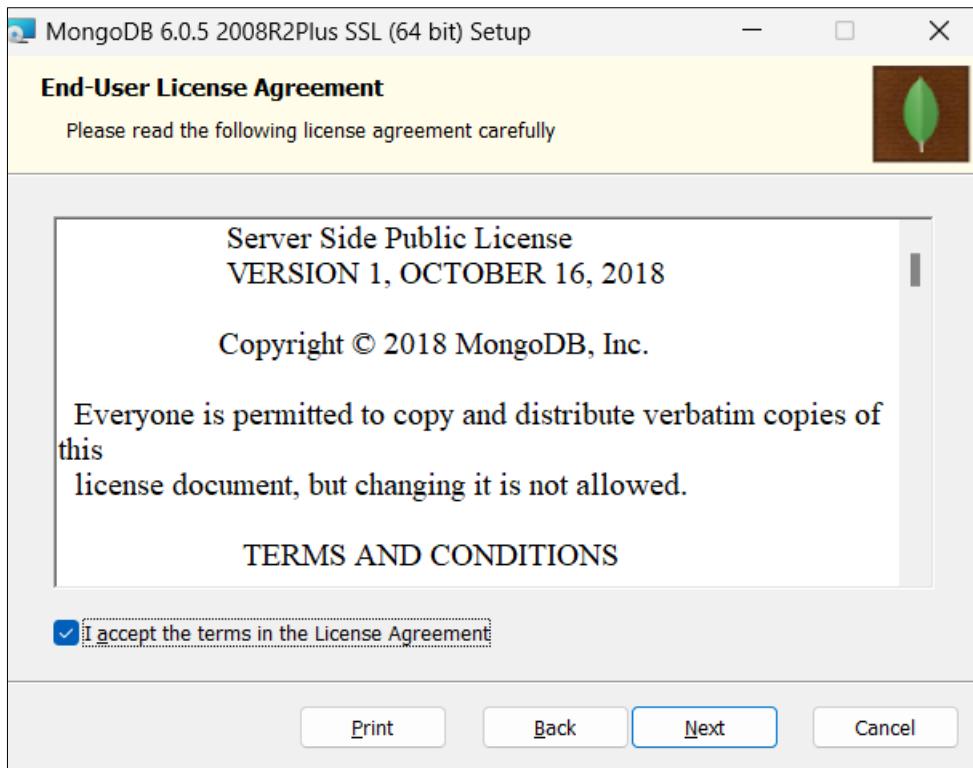
The MongoDB 6.0.5 2008R2Plus SSL (64-bit) Setup wizard opens as shown in Figure 1.1.



**Figure 1.1: Welcome Page of the MongoDB Installation Wizard**

5. Click **Next**.

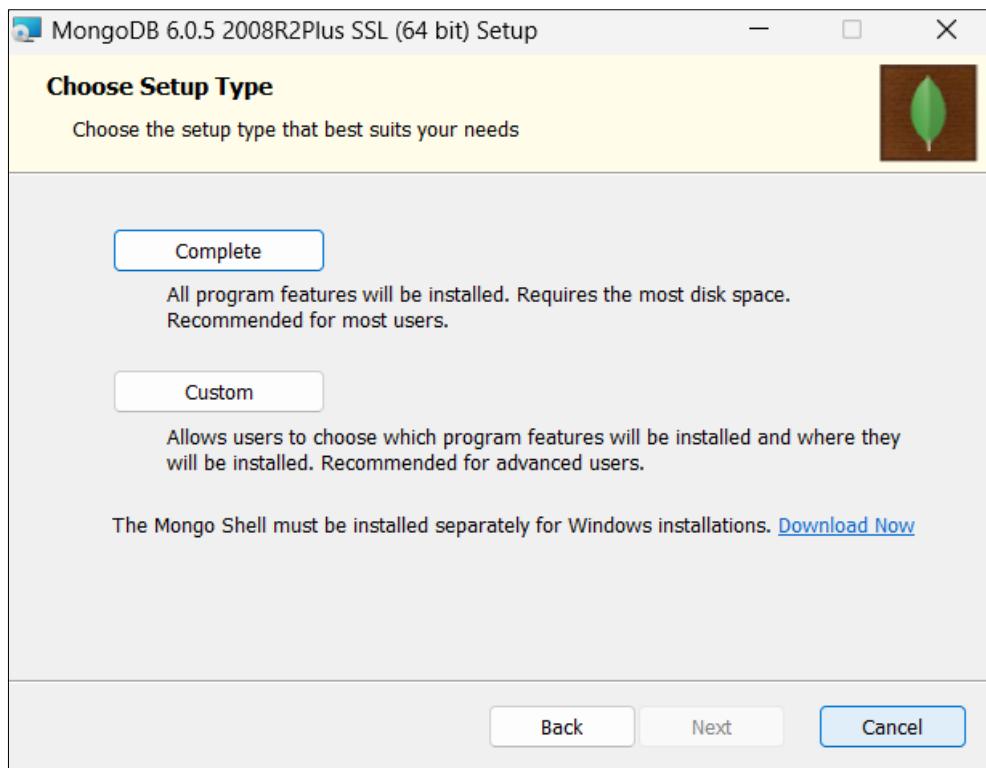
The End-User License Agreement page of the wizard opens as shown in Figure 1.2.



**Figure 1.2: End-User License Agreement Page**

6. To proceed, select the **I accept the terms in the License Agreement** check box.
7. Click **Next**.

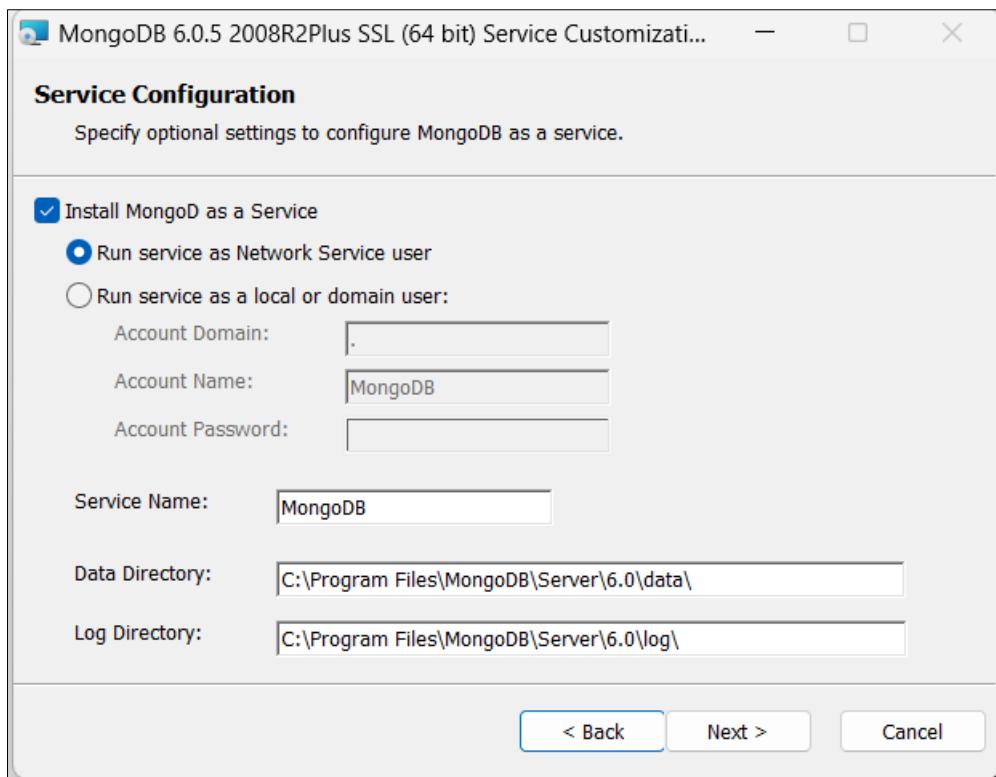
The Choose Setup Type page of the wizard opens as shown in Figure 1.3.



**Figure1.3: Choose Setup Type Page**

8. Click **Complete**.

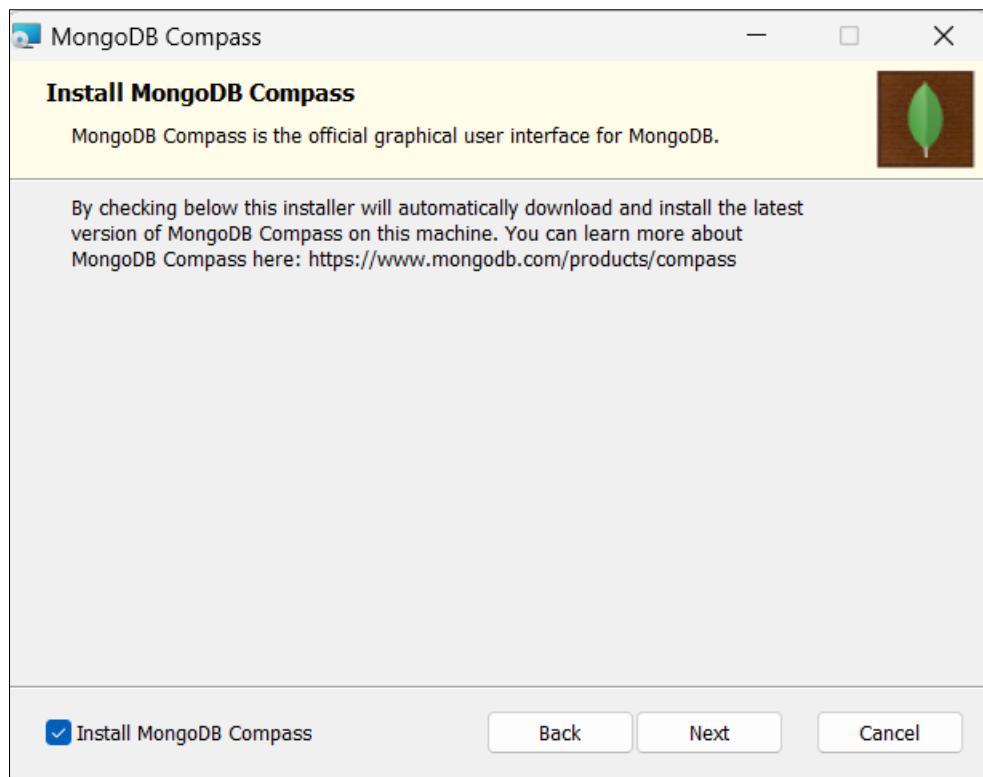
The Service Configuration page of the wizard opens as shown in Figure 1.4.



**Figure 1.4: Service Configuration Page**

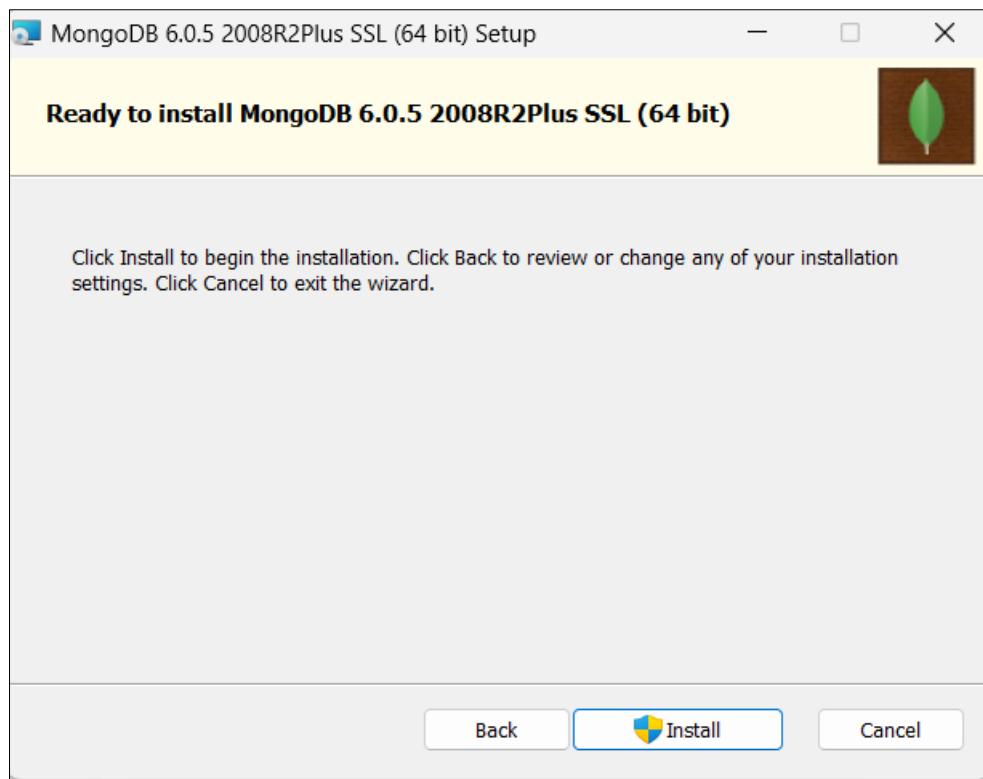
9. Click **Next**.

The Install MongoDB Compass page of the wizard opens as shown in Figure 1.5.



**Figure 1.5: Install MongoDB Compass Page**

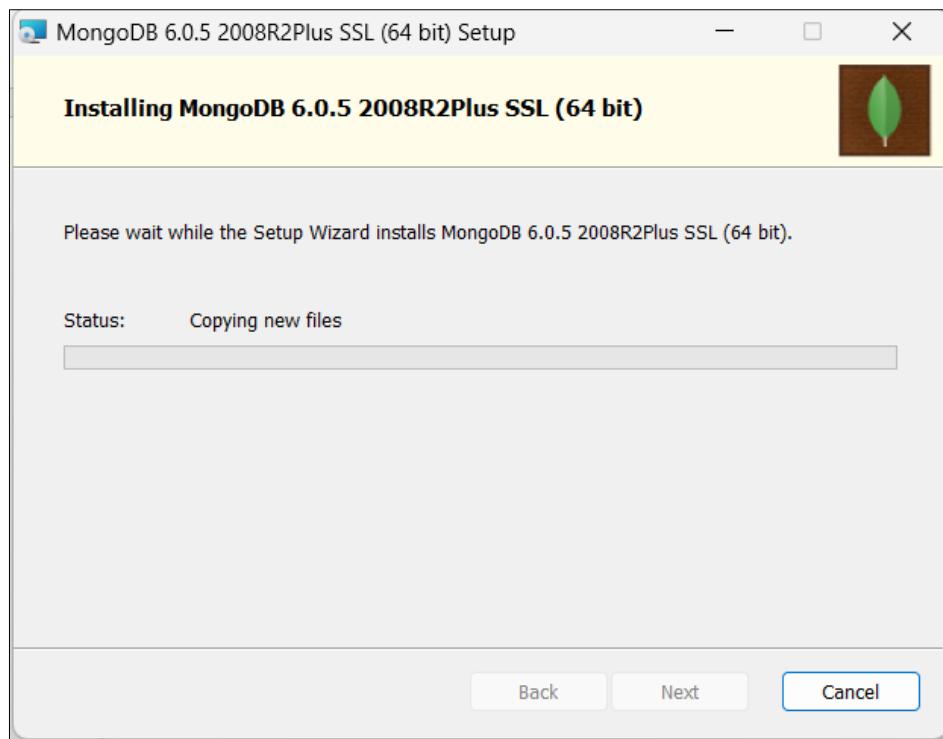
10. Ensure that the **Install MongoDB Compass** check box is selected.
11. Click **Next**.  
The Ready to install MongoDB 6.0.5 2008R2Plus SSL (64-bit) page of the wizard opens as shown in Figure 1.6.



**Figure 1.6: Screenshot of the Ready to Install Page**

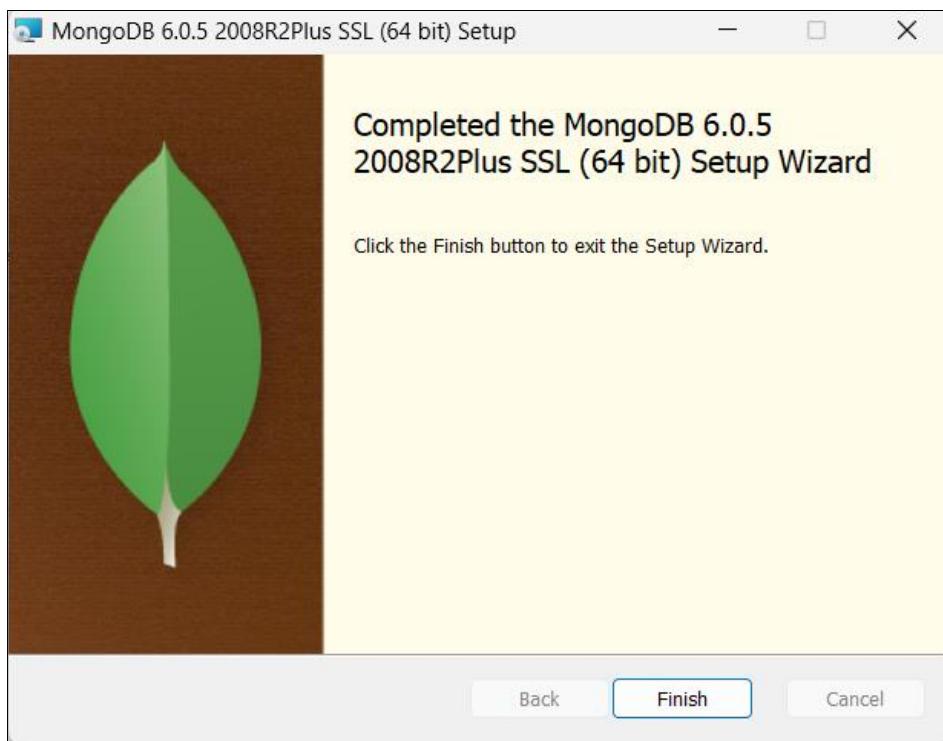
12. Click **Install**.

The Installing MongoDB 6.0.5 2008R2Plus SSL (64-bit) page of the wizard opens as shown in Figure 1.7.



**Figure 1.7: Installation in Progress Page**

The Completed the MongoDB 6.0.5 2008R2Plus SSL (64-bit) Setup Wizard page of the wizard opens as shown in Figure 1.8.



**Figure 1.8: Installation Completion Page**

13. Click **Finish**.



MongoDB Compass is also installed along with MongoDB Community Server. The MongoDB Compass window opens by default. Close this window.

## 1.4 Installing MongoDB Shell

After installing the MongoDB Server, install the MongoDB Shell.

To install the MongoDB Shell:

1. Open a browser and navigate to the URL:

<https://www.mongodb.com/try/download/shell>

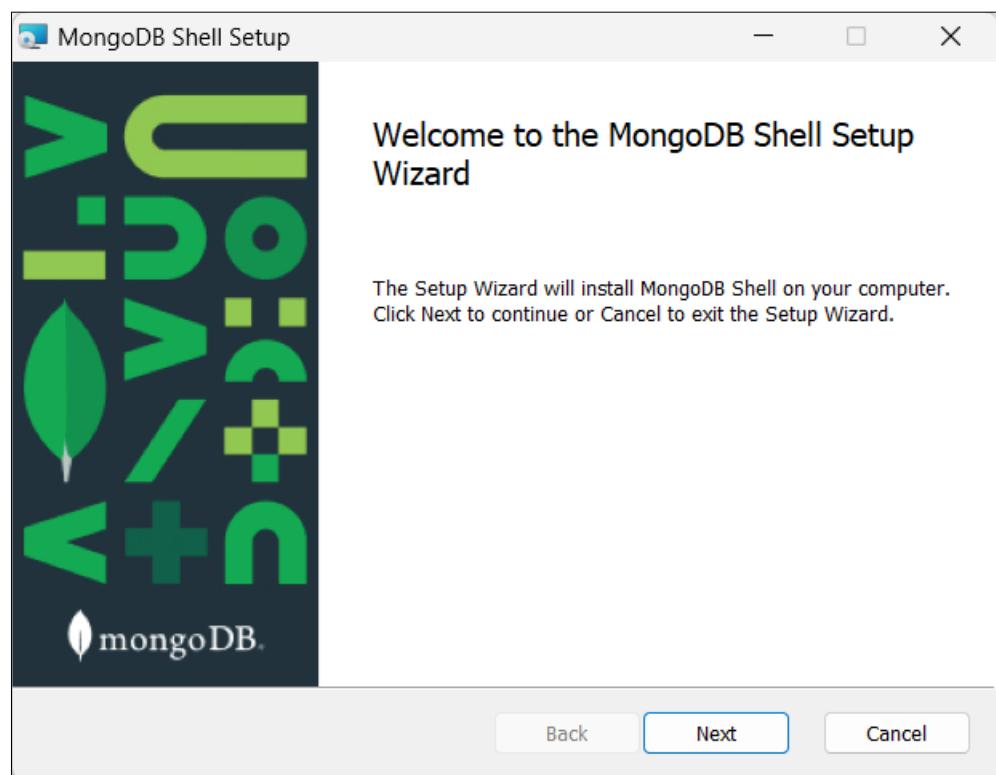
The MongoDB Shell Download page opens as shown in Figure 1.9.

The screenshot shows a web-based download interface for MongoDB Shell. It has three main dropdown menus: 'Version' set to '1.8.2', 'Platform' set to 'Windows 64-bit (8.1+) (MSI)', and 'Package' set to 'msi'. Below these dropdowns is a large search bar containing the text 'msi'.

**Figure 1.9: Version, Platform, and Package Selection**

2. On this page, scroll down and ensure that:
  - In the **Version** drop-down, **1.8.2** is selected.
  - In the **Platform** drop-down, **Windows 64-bit (8.1+) (MSI)** is selected.
  - In the **Package** drop-down, **msi** is selected.
3. Click **Download**.
4. After the download is complete, run the installer.

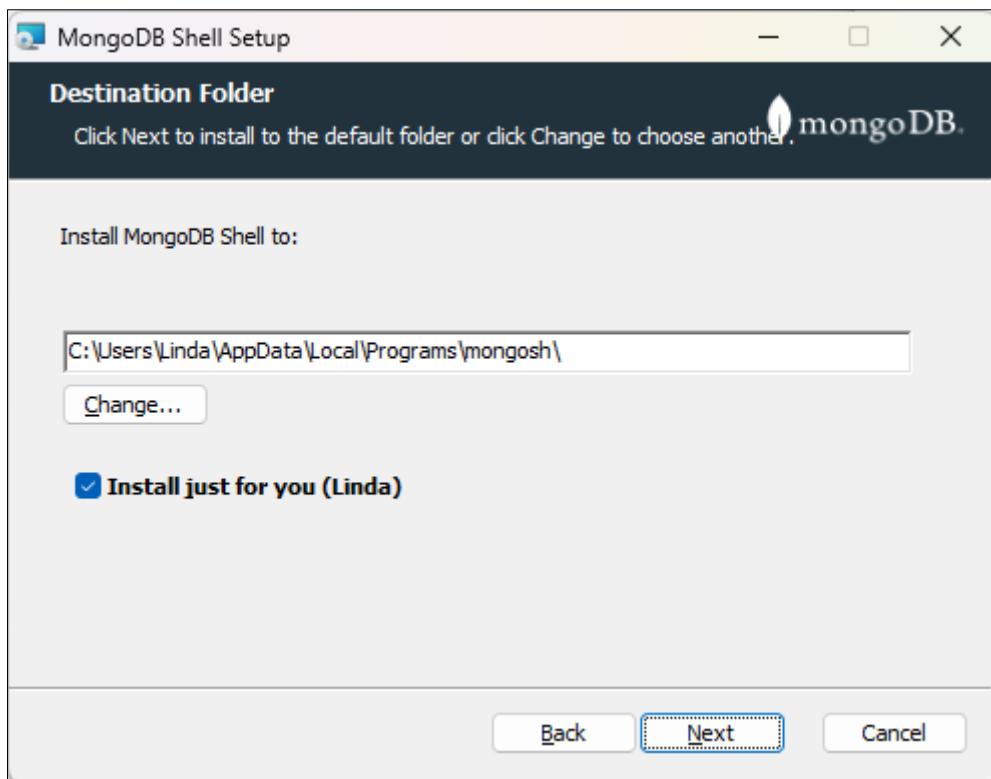
The MongoDB Shell Setup wizard opens as shown in Figure 1.10.



**Figure 1.10: Screenshot of MongoDB Shell Setup Wizard**

5. Click **Next**.

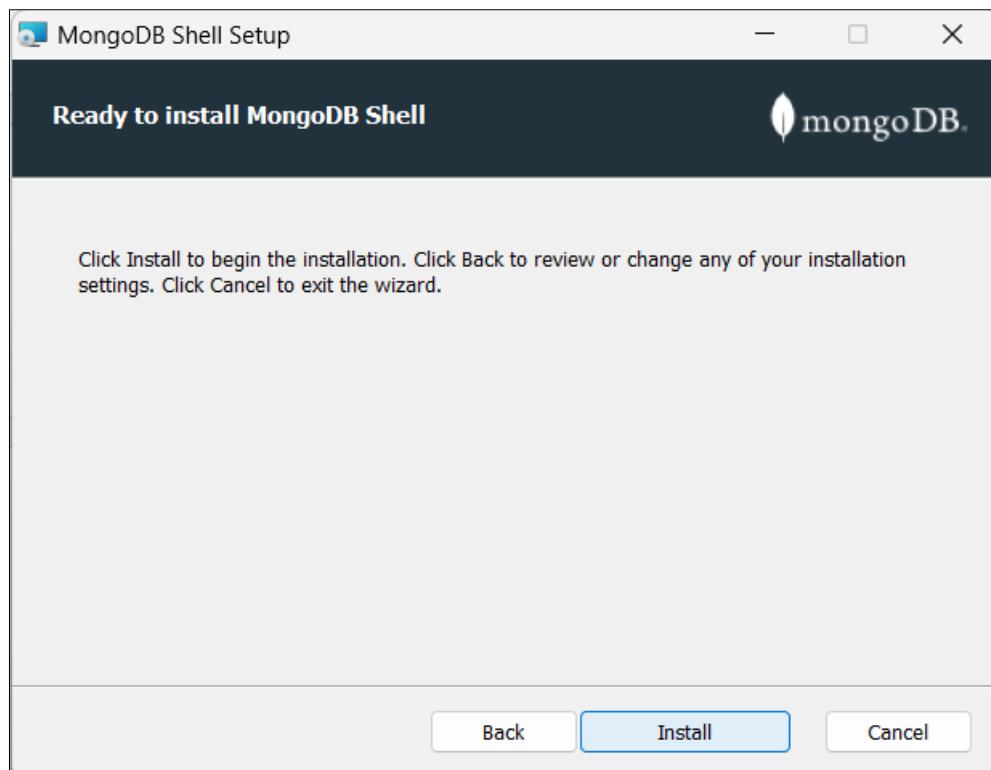
The Destination Folder page of the wizard opens as shown in Figure 1.11.



**Figure 1.11: Destination Folder Page**

6. Check the folder path and click **Next**.

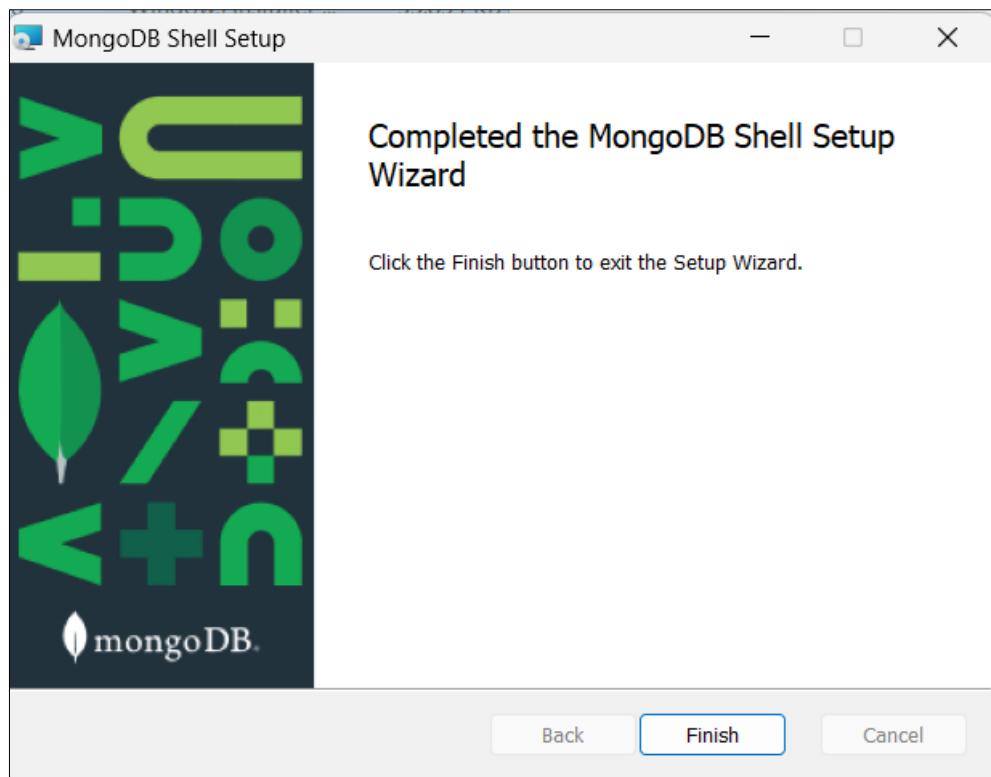
The Ready to install MongoDB Shell page of the wizard opens as shown in Figure 1.12.



**Figure 1.12: Ready to Install MongoDB Shell Page**

7. Click **Install**.

The MongoDB Shell Setup page of the wizard opens as shown in Figure 1.13.



**Figure 1.13: Screenshot of the MongoDB Shell Setup Page**

8. Click **Finish**.

## **1.5 Setting up the Environment**

To set up the environment:

1. Open a Windows command prompt (cmd.exe) as an **Administrator**.
2. Create the data directory where MongoDB will store all the data, as shown in Figure 1.14.

```
C:\>md "\data\db"  
C:\>
```

**Figure 1.14: Creating Directories**



MongoDB listens for connections from clients only on port 27017, by default. The default storage path for data is /data/db directory.

To start the **MongoDB** database, run the command:

```
"C:\Program Files\MongoDB\Server\6.0\bin\mongod.exe" --dbpath="c:\data\db"
```

The command is executed as shown in Figure 1.15.

```
C:\>"C:\Program Files\MongoDB\Server\6.0\bin\mongod.exe" --dbpath="c:\data\db"
[{"t": {"$date": "2023-05-01T13:09:48.839+05:30"}, "s": "I", "c": "NETWORK", "id": 4915701, "ctx": "thread1", "msg": "Initialized wire specification", "attr": {"spec": {"incomingExternalClient": {"minWireVersion": 0, "maxWireVersion": 17}, "incomingInternalClient": {"minWireVersion": 0, "maxWireVersion": 17}, "outgoing": {"minWireVersion": 6, "maxWireVersion": 17}, "isInternalClient": true}}}, {"t": {"$date": "2023-05-01T13:09:50.049+05:30"}, "s": "I", "c": "CONTROL", "id": 23285, "ctx": "thread1", "msg": "Automatically disabling TLS 1.0, to force-enable TLS 1.0 specify --sslDisabledProtocols 'none'"}, {"t": {"$date": "2023-05-01T13:09:50.053+05:30"}, "s": "I", "c": "NETWORK", "id": 4648602, "ctx": "thread1", "msg": "Implicit TCP FastOpen in use."}, {"t": {"$date": "2023-05-01T13:09:50.069+05:30"}, "s": "I", "c": "REPL", "id": 5123008, "ctx": "thread1", "msg": "Successfully registered PrimaryOnlyService", "attr": {"service": "TenantMigrationDonorService", "namespace": "config.tenantMigrationDonors"}}, {"t": {"$date": "2023-05-01T13:09:50.069+05:30"}, "s": "I", "c": "REPL", "id": 5123008, "ctx": "thread1", "msg": "Successfully registered PrimaryOnlyService", "attr": {"service": "TenantMigrationRecipientService", "namespace": "config.tenantMigrationRecipients"}}
```

**Figure 1.15: Starting the MongoDB Database**

The --dbpath option shows the database directory. If the MongoDB database server is running correctly, the message: **waiting for connections** is displayed as shown in Figure 1.16.

```
[{"t": {"$date": "2023-05-15T17:39:03.517+05:30"}, "s": "I", "c": "NETWORK", "id": 23015, "ctx": "listener", "msg": "Listening on", "attr": {"address": "127.0.0.1"}}, {"t": {"$date": "2023-05-15T17:39:03.517+05:30"}, "s": "I", "c": "NETWORK", "id": 23016, "ctx": "listener", "msg": "Waiting for connections", "attr": {"port": 27017, "ssl": "off"}}]
```

**Figure 1.16: Waiting for Connections Message**

Now, the MongoDB instance mongod is successfully running.



Add C:\Program Files\MongoDB\Server\6.0\bin to the system path. This allows to run the mongod instance instantly without having to navigate to this folder each time and then run the instance.

## 1.6 Connecting MongoDB Server and MongoDB Shell

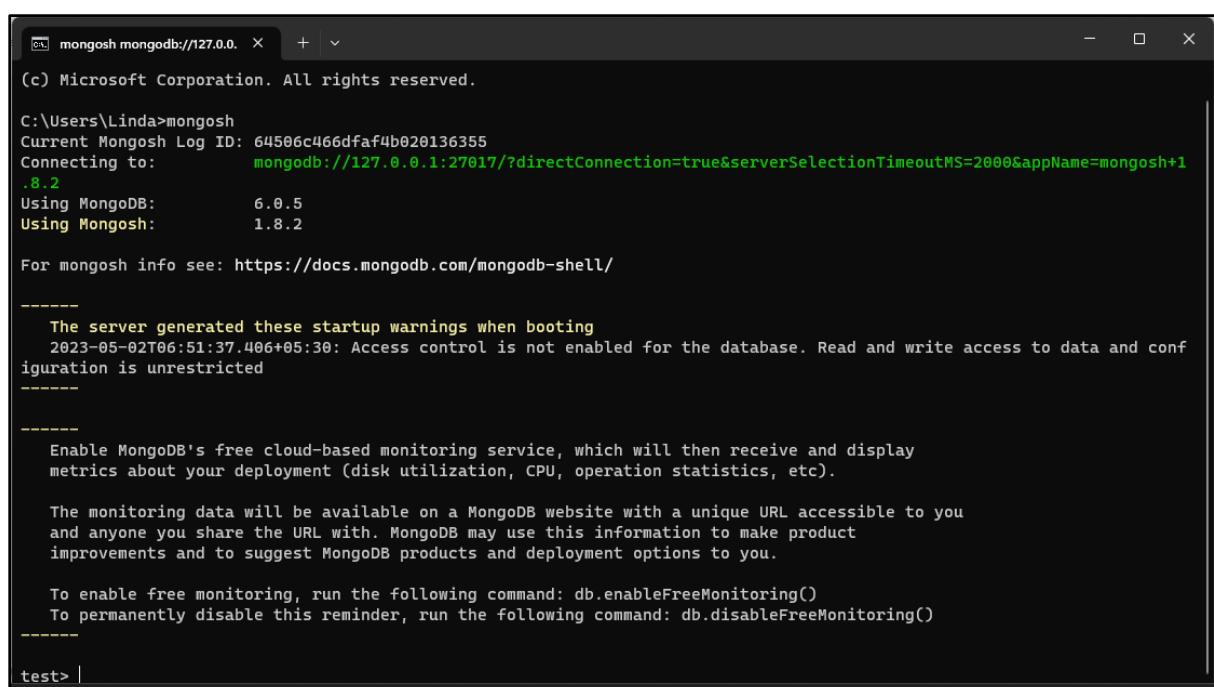
---

To connect the MongoDB Shell with MongoDB Server:

1. Open another Windows command prompt/interpreter.
2. To connect the MongoDB Shell to a MongoDB deployment running on **localhost** with the default port 27017, run the command:

```
mongosh
```

The command is executed as shown in Figure 1.17.



A screenshot of a Windows command prompt window titled "mongosh mongodb://127.0.0.1:27017/?directConnection=true&serverSelectionTimeoutMS=2000&appName=mongosh+1.8.2". The window shows the following text:

```
(c) Microsoft Corporation. All rights reserved.

C:\Users\Linda>mongosh
Current Mongosh Log ID: 64506c466dfaf4b020136355
Connecting to:      mongodb://127.0.0.1:27017/?directConnection=true&serverSelectionTimeoutMS=2000&appName=mongosh+1.8.2
Using MongoDB:      6.0.5
Using Mongosh:     1.8.2

For mongosh info see: https://docs.mongodb.com/mongodb-shell/

-----
The server generated these startup warnings when booting
2023-05-02T06:51:37.406+05:30: Access control is not enabled for the database. Read and write access to data and configuration is unrestricted
-----

-----
Enable MongoDB's free cloud-based monitoring service, which will then receive and display
metrics about your deployment (disk utilization, CPU, operation statistics, etc).

The monitoring data will be available on a MongoDB website with a unique URL accessible to you
and anyone you share the URL with. MongoDB may use this information to make product
improvements and to suggest MongoDB products and deployment options to you.

To enable free monitoring, run the following command: db.enableFreeMonitoring()
To permanently disable this reminder, run the following command: db.disableFreeMonitoring()
-----
```

test> |

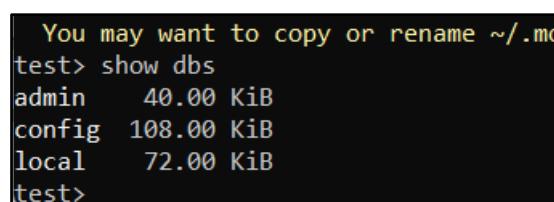
**Figure 1.17: Output of the Mongosh Command**

MongoDB Shell is now successfully connected to the MongoDB server.

3. To check the connection, run the command:

```
show dbs
```

List of existing databases is displayed as shown in Figure 1.18



```
You may want to copy or rename ~/.mc
test> show dbs
admin   40.00 KiB
config  108.00 KiB
local   72.00 KiB
test>
```

**Figure 1.18: Output of the show dbs Command**

## **1.7 Summary**

---

- A NoSQL database is a non-relational data that can support huge volumes of unstructured data.
- MongoDB is an open-source NoSQL database that stores data in the BSON format.
- The drivers, shell, and storage engine are vital elements for MongoDB to function.
- MongoDB stores data in the form of databases, collections, and documents.
- MongoDB works in two layers: application layer and data layer.
- Mongo DB Compass is the GUI for interacting with the MongoDB Server.
- MongoDB Atlas is the product MongoDB offers to compute on the cloud.
- To set up the MongoDB environment, first, install the MongoDB Server, then install the MongoDB Shell, and then use mongosh to connect the server to the shell.

## **Test Your Knowledge**

---

1. Which of the following options is the data storage model of NoSQL databases?
  - a. Document
  - b. Tables with fixed rows and columns
  - c. Key-value
  - d. Graph
  
2. Which of the following statements are not true about NoSQL databases?
  - a. It has flexible schema.
  - b. It supports Horizontal scaling.
  - c. The data architectures offered by NoSQL databases are more versatile than the data architectures offered by the relational database tables.
  - d. It offers only read scalability.
  
3. Document database stores data in which of the following format?
  - a. JSON
  - b. BSON
  - c. Google Docs
  - d. XML
  
4. Which of the following is the default port for MongoDB which listens to the client for connection?
  - a. 27017
  - b. 20717
  - c. 21071
  - d. 20171
  
5. Which of the following service does MongoDB cloud Atlas provides?
  - a. Infrastructure as a service
  - b. Platform as a service
  - c. Data as a service
  - d. Function as a service

## **Answers to Test Your Knowledge**

---

1	a, c, d
2	d
3	a, b, d
4	a
5	c

## Try it Yourself

---

1. Install the **MongoDB 6.0** Community edition on your Windows.
2. Install **Mongo Shell 1.8.2** on your Windows system.
3. Using command prompt, set up the environment for MongoDB to execute `mongod` instance and connect `mongosh` with `mongod` instance. Execute a command to view the databases.



## SESSION 2

### MONGODB DATABASES

---

#### Learning Objectives

In this session, students will learn to:

- Explain the installation of MongoDB Database Tools and set the environment variable for MongoDB Database Tools
- Explain the method to load a sample dataset into MongoDB server
- Describe different datatypes in MongoDB
- Explain the methods to create database and collection
- Describe the ways to insert, query, update, and delete documents, collections, and databases

This session will provide an overview of installing MongoDB Database Tools and setting the environment variables for it. It will also explain how to load a sample database into MongoDB server. This session will explore the methods provided by MongoDB to create databases and collections. It will also explain the methods to insert, query, update, and delete documents, collections, and databases.

#### 2.1 Install MongoDB Database Tools

MongoDB provides a collection of tools that facilitate the users to work with and manage databases. Table 2.1 describes some of the common MongoDB

tools.

Tool Name	Description
mongodump	Exports and creates a binary backup of the contents of a database
mongorestore	Restores backup data from mongodump into a database and loads standard input data into a mongod or mongos instance
mongoexport	Exports and creates a JavaScript Object Notation (JSON) or Comma-Separated Values (CSV) backup of the contents stored in a MongoDB instance
mongoimport	Imports data from CSV, JSON, and Tab-Separated Values (TSV) files into a MongoDB instance
bsondump	Converts BSON(Binary JavaScript Object Notation) files into JSON files
mongotop	Provides data on the amount of time spent by a MongoDB instance on reading and writing
mongofiles	Facilitates the manipulation of files stored in a MongoDB instance
mongostat	Provides statistics such as the number of inserts, deletes, and update queries executed per second and the amount of virtual memory used by the process

**Table 2.1: Some Common MongoDB Tools**

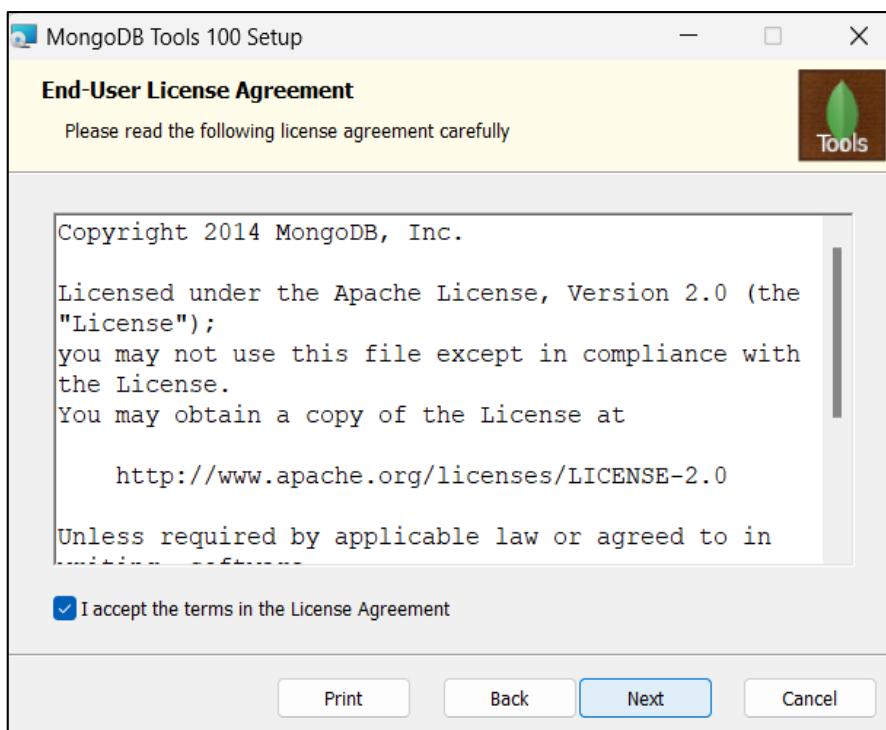
To install MongoDB database tools for Windows:

1. Open the browser and navigate to the Website  
<https://www.mongodb.com/try/download/database-tools>
2. On this page, scroll down:
  - In the **Version** drop-down, ensure that **100.7.0** is selected.
  - In the **Platform** drop-down, ensure that **Windows x86\_64** is selected.
  - In the **Package** drop-down, select the **msi** option.
3. Click **Download**.
4. After the download is complete, run the installer.  
The **MongoDB Tools 100 Setup** wizard opens as shown in Figure 2.1.



**Figure 2.1: MongoDB Tools Setup wizard**

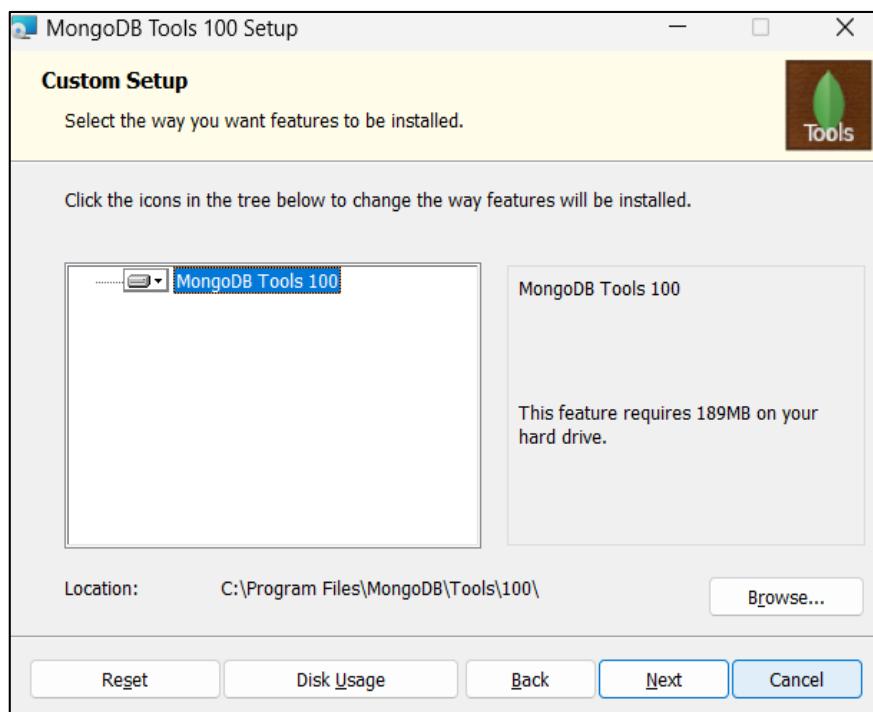
5. Click **Next**.  
The **End-User License Agreement** page opens.
6. Select the **I accept the terms in the License Agreement** check box as shown in Figure 2.2.



**Figure 2.2: End-User License Agreement**

7. Click **Next**.

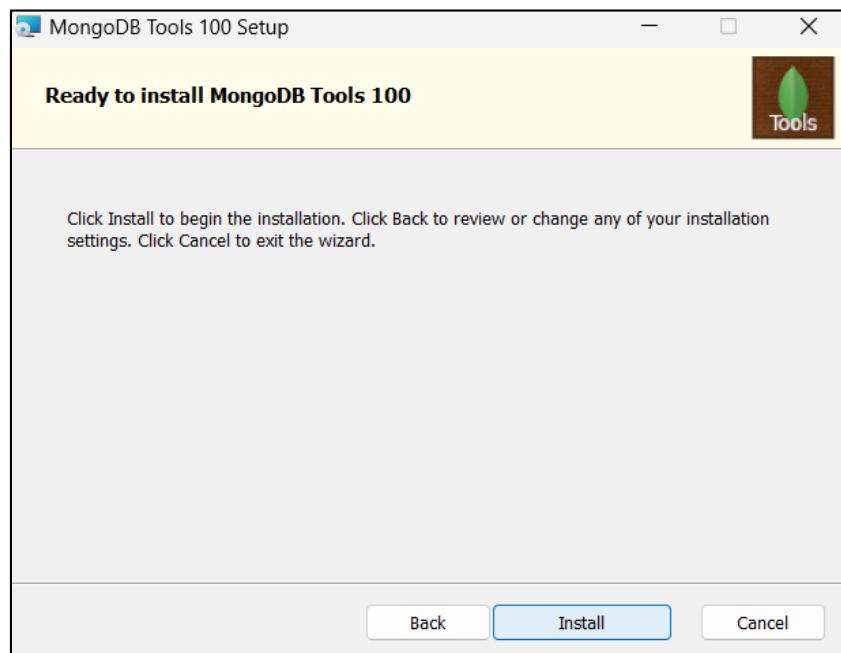
The **Custom Setup** page opens as shown in Figure 2.3.



**Figure 2.3: Custom Setup Page**

8. Click **Next**.

The **Ready to Install MongoDB Tools 100** page opens as shown in Figure 2.4.



**Figure 2.4: Ready to Install MongoDB Tools 100**

9. Click **Install**.

The **Completed the MongoDB Tools 100 Setup Wizard** page opens as shown in Figure 2.5.

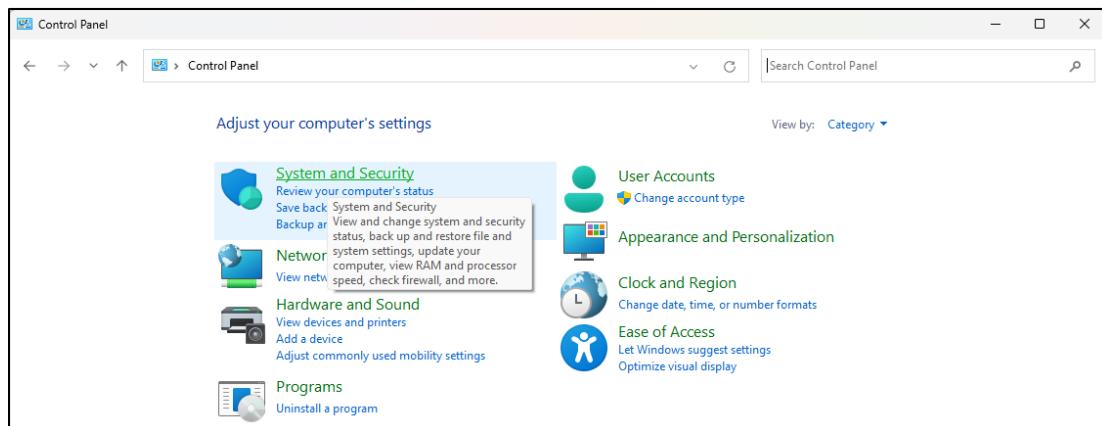


**Figure 2.5: Completed the MongoDB Tools 100 Setup Wizard**

10. Click **Finish**.

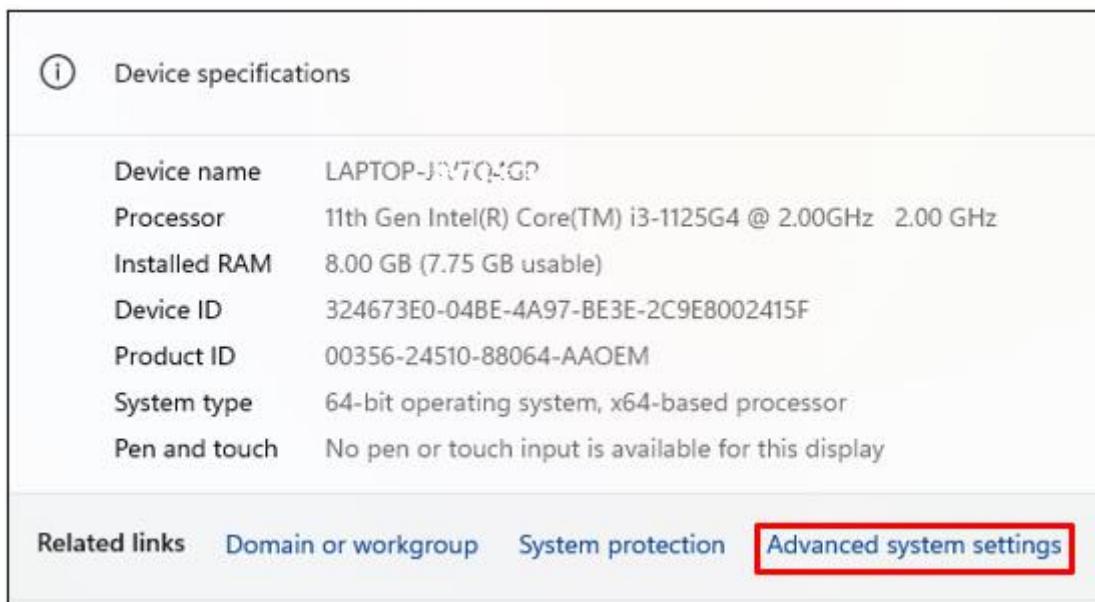
Now, MongoDB Database Tools are installed. The next step is to set the environment variables for the executable file. To set the environment variable:

1. Open the **Control Panel** in **Windows**.
2. In the **Control Panel** window, select **System and Security** as shown in Figure 2.6.



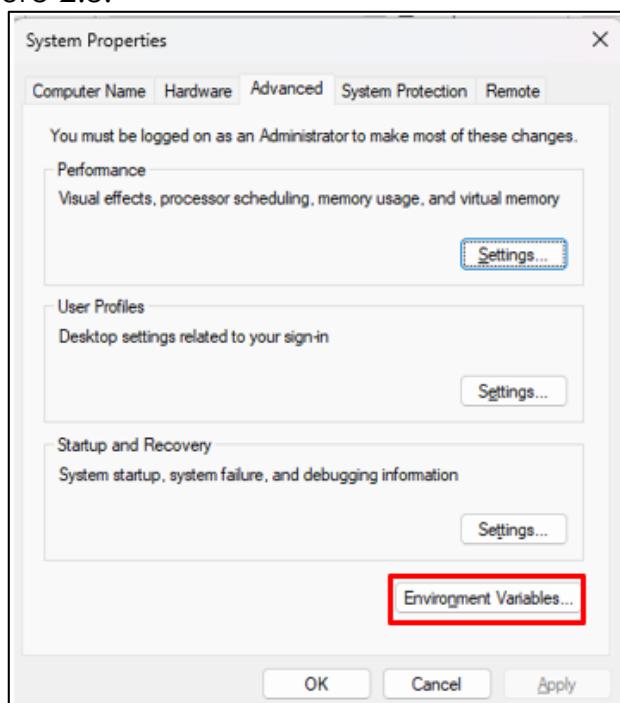
**Figure 2.6: Control Panel**

3. In the **System and Security** window that opens, click **System**.
4. In the **System → About** window that opens, select **Advanced system settings** as shown in Figure 2.7.



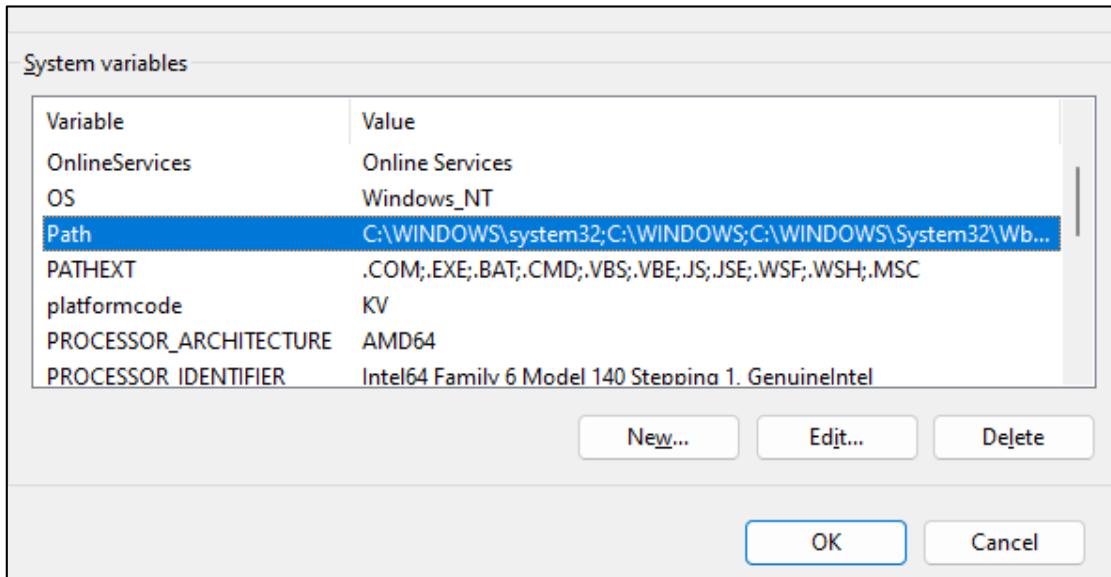
**Figure 2.7: Advanced System Settings**

- The **System Properties** window opens.
5. In the **System Properties** window, select **Environment Variables** as shown in Figure 2.8.



**Figure 2.8: System Properties**

6. In the **Environment Variables** window that opens, under **System variables** section, select **Path** and click **Edit** as shown in Figure 2.9.



**Figure 2.9: Path**

7. In the **Edit environment variable** window that opens, click **New**, and then click **Browse** to navigate to the folder where the database tools are installed.  
 8. Click **OK** in all the windows opened in Steps 5 through 7.  
 The environment variable is now set for MongoDB tools.

## 2.2 Load Sample Dataset into MongoDB Server

---

To load a sample dataset into MongoDB Server using the `mongoimport` tool:

1. Copy the sample dataset named `Sample_dataset` to the C drive.
2. Open the command prompt.
3. To run the `mongod` instance, execute the command:

```
"C:\Program Files\MongoDB\Server\6.0\bin\mongod.exe" --dbpath="c:\data\db"
```

The command is executed as shown in Figure 2.10.

```
C:\>"C:\Program Files\MongoDB\Server\6.0\bin\mongod.exe" --dbpath="c:\data\db"
{"t": {"$date": "2023-05-15T17:39:00.769+05:30"}, "s": "I", "c": "CONTROL", "id": 23285, "ctx": "thread1", "msg": "Automatically disabling TLS 1.0, to force-enable TLS 1.0 specify --sslDisabledProtocols 'none'"}
{"t": {"$date": "2023-05-15T17:39:02.257+05:30"}, "s": "I", "c": "NETWORK", "id": 4915701, "ctx": "thread1", "msg": "Initialized wire specification", "attr": {"spec": {"incomingExternalClient": {"minWireVersion": 0, "maxWireVersion": 17}, "incomingInternalClient": {"minWireVersion": 0, "maxWireVersion": 17}, "outgoing": {"minWireVersion": 6, "maxWireVersion": 17}, "isInternalClient": true}}}
```

**Figure 2.10: Starting the MongoDB Database**

The `--dbpath` option shows the database directory. If the MongoDB database server is running correctly, the message waiting for connections is displayed as shown in Figure 2.11.

```
{"t": {"$date": "2023-05-15T17:39:03.517+05:30"}, "s": "I", "c": "NETWORK", "id": 23015, "ctx": "listener", "msg": "Listener is listening on", "attr": {"address": "127.0.0.1"}}, {"t": {"$date": "2023-05-15T17:39:03.517+05:30"}, "s": "I", "c": "NETWORK", "id": 23016, "ctx": "listener", "msg": "Waiting for connections", "attr": {"port": 27017, "ssl": "off"}}
```

**Figure 2.11: Waiting for Connections Message**

Now, the MongoDB instance `mongod` is successfully running.

4. Open a new command prompt and type the command as:

```
mongoimport --db sample_weatherdata --collection data --jsonArray --file C:\Sample_dataset\sample_weatherdata\data.json
```

The command executes as shown in Figure 2.12.

```
PS C:\Users\linda> mongoimport --db sample_weatherdata --collection data --jsonArray --file C:\Sample_dataset\sample_weatherdata\data.json
2023-05-15T17:54:39.729+0530      connected to: mongodb://localhost/
2023-05-15T17:54:40.578+0530      10000 document(s) imported successfully. 0 document(s) failed to import.
```

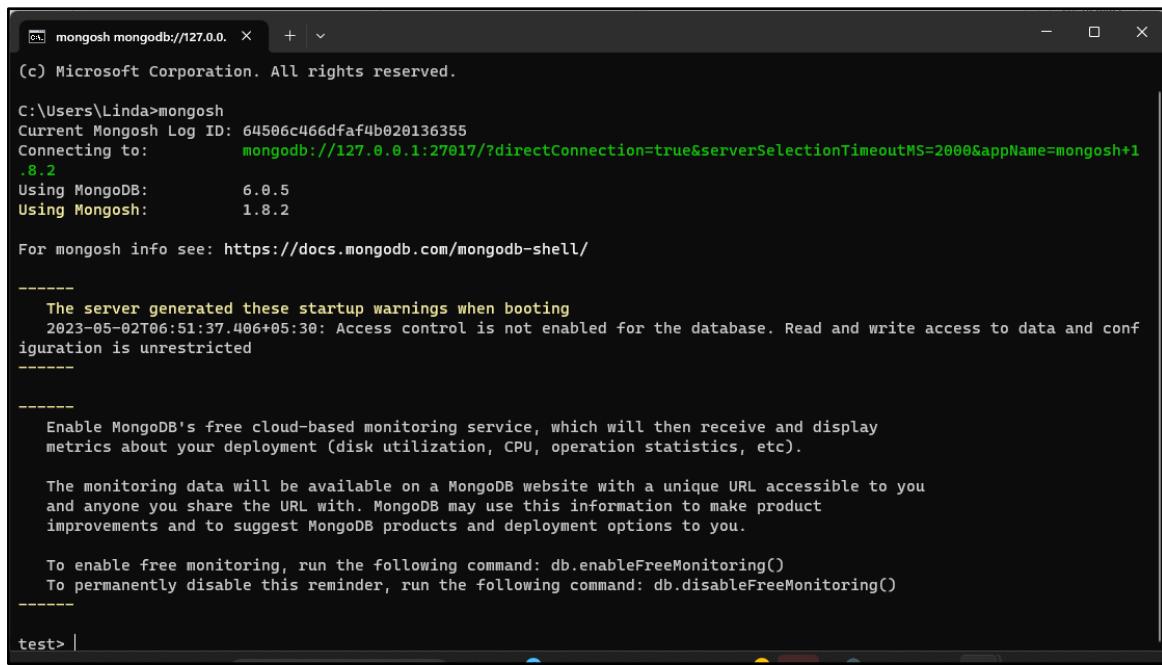
**Figure 2.12: Import Sample Data**

The sample dataset is now imported into the `mongod` instance.

5. To verify this, run the MongoDB shell command as:

```
mongosh
```

The command executes as shown in Figure 2.13.



```
[C:\mongosh mongodb://127.0.0.1:27017/?directConnection=true&serverSelectionTimeoutMS=2000&appName=mongosh+1.8.2] mongosh
(c) Microsoft Corporation. All rights reserved.

C:\Users\Linda>mongosh
Current MongoDB Log ID: 64506c466dfaf4b020136355
Connecting to:      mongodb://127.0.0.1:27017/?directConnection=true&serverSelectionTimeoutMS=2000&appName=mongosh+1.8.2
Using MongoDB:      6.0.5
Using Mongosh:      1.8.2

For mongosh info see: https://docs.mongodb.com/mongodb-shell/

-----
The server generated these startup warnings when booting
2023-05-02T06:51:37.406+05:30: Access control is not enabled for the database. Read and write access to data and configuration is unrestricted
-----

Enable MongoDB's free cloud-based monitoring service, which will then receive and display metrics about your deployment (disk utilization, CPU, operation statistics, etc).

The monitoring data will be available on a MongoDB website with a unique URL accessible to you and anyone you share the URL with. MongoDB may use this information to make product improvements and to suggest MongoDB products and deployment options to you.

To enable free monitoring, run the following command: db.enableFreeMonitoring()
To permanently disable this reminder, run the following command: db.disableFreeMonitoring()
-----
test> |
```

**Figure 2.13: mongosh Command**

6. To view the databases available in this mongod instance, execute the command as:

```
show dbs
```

The command executes as shown in Figure 2.14.

```
test> show dbs
admin              40.00 KiB
config             60.00 KiB
local              72.00 KiB
sample_weatherdata  2.68 MiB
```

**Figure 2.14: show dbs Command**

This confirms that the sample data named `sample_weatherdata` is imported into the mongod instance.



---

The `show dbs` command would display a database only when the database contains at least one document.

---

## 2.3 Datatypes in MongoDB

---

MongoDB stores data internally in BSON format. BSON is an extension of JSON. It supports all the data types that JSON supports. Additionally, it supports other datatypes such as date, timestamp, and binary data. Table 2.2 describes the

datatypes that are available in MongoDB.

Datatype	Description	Example
String	A String is used to store text in MongoDB. BSON strings are of Universal Coded Character Set (UCS) Transformation Format (UTF – 8).	Store the name of a student or employee, designation of an employee, and address
Integer	An Integer is used to store numeric values. MongoDB offers two variations of integer datatype—32-bit signed integer and 64-bit signed integer.	Store numeric values such as marks of a student, salary of an employee, and count of products in inventory
Double	A Double is used to store numeric values that contain 8-byte floating point values.	Store temperature, revenue, and price of an item
Decimal	Decimal is used to store 128-bit floating point decimal values.	Store financial and scientific computations
Boolean	Boolean is used to store true or false as values for the fields.	Store the value of 'Is_married', 'Is_fulltime_employee', and 'Is_passed'
Null	Null is used to store NULL value.	Store the driver's license number of people who do not own a driver's license
Array	Array is a collection of multiple values that are of same type or different types.	Store marks of a student in five subjects
Object	Object is used to store embedded documents.	Store the description of an item which would include the item's shape, color, and dimensions
ObjectID	ObjectID is a unique identifier for each document in a collection. It is a hexadecimal string of 12-bytes that includes: <ul style="list-style-type: none"> <li>• Timestamp value (4-bytes)</li> <li>• Random value(5-bytes with 3-bytes of machine ID and 2-bytes of process ID)</li> <li>• Counter (3-bytes)</li> </ul>	A hexadecimal ID that MongoDB assigns to each new field in MongoDB, for example, '64617a0420cf1016275bd88f'
Date	Stores date as a 64-bit signed integer which is a Universal Time Co-ordinated (UTC)	String date: Mon May 15 2023 01:03:25 GMT-0500 (EST)

Datatype	Description	Example
	<p>datetime format. This BSON date format allows storage of dates before and after Jan 1, 1970, 00:00:00 using the negative and positive sign-bits, respectively.</p> <p>When the user creates a date variable in MongoDB, it returns the date in string format.</p> <p>When the user creates a date variable in MongoDB using the new keyword, MongoDB returns it in the International Standards Organization Date (ISODate) format.</p>	<p>ISODate: 2023-05-15T04:25:48.512Z</p>

**Table 2.2: Datatypes in MongoDB**

## 2.4 Working with Databases and Collections

Consider that the user wants to create a database in MongoDB and perform certain operations in it. MongoDB offers a set of commands that facilitates the user to create and drop databases and collections. MongoDB also provides commands that can be used to insert, query, update, and delete documents in a collection.

### 2.4.1 Create Database

Consider that the user wants to create a database named `Student_detail`. The syntax to create a database is:

```
use DATABASE_NAME
```

To create a database named `Student_detail`:

1. Ensure that the `mongod` instance is running.
2. To run the MongoDB Shell, at the command prompt, execute the command as:

```
mongosh
```

3. To create a database named `Student_detail`, run the command as:

```
use Student_detail
```

The command executes as shown in Figure 2.15.

```
test> use Student_detail
switched to db Student_detail
Student_detail> |
```

**Figure 2.15: Create Database**

The database named `Student_detail` has been created.

#### 2.4.2 Create Collection

Consider that the user wants to create a collection named `Studentinfo` in the `Student_detail` database. The syntax to create a collection or a view is:

```
db.createCollection(name, options)
```

Table 2.3 describes the parameters of the `createCollection` command.

Parameter	Type	Description
name	String	This parameter specifies the name of the collection to be created. This is a mandatory parameter.
options	Document	<p>This parameter is optional. It includes fields that specify whether the entity to be created must be a:</p> <ul style="list-style-type: none"> <li>• <b>Capped collection:</b> Capped collections are collections of fixed size. If the maximum size of the collection is reached, then the oldest documents are erased to create space for new entries. The size parameter and the maximum number of documents that can be created must be specified for this option.</li> <li>• <b>Timeseries collection:</b> Timeseries collections are collections that store sequences of data points (key-value pairs) over a period along with the time at which the data points were recorded.</li> <li>• <b>Clustered collection:</b> Clustered collections are collections that are stored in the order of a clustered index—a key value that specifies the order of arrangement of the documents in a collection. The parameter includes fields such as key value field, uniqueness of the clustered index, and name of the index.</li> <li>• <b>View:</b> View is a virtual collection. If a view must be created, the name of the source collection or view must be specified.</li> </ul> <p>This parameter also includes a field named <code>validator</code> that allows the users to create collections with user-defined validation rules.</p>

**Table 2.3: Parameters of `createCollection()` Method**

To create a collection named `Studentinfo` in the database named `Student_detail`:

1. To switch to the `Student_detail` database, execute the command as:

```
use Student_detail
```

2. To create a collection named `Studentinfo`, execute the command

as:

```
db.createCollection("Studentinfo")
```

The command executes as shown in Figure 2.16.

```
Student_detail> db.createCollection("Studentinfo")
{ ok: 1 }
Student_detail> |
```

**Figure 2.16: Create Collection**

The Studentinfo collection is created.

3. To view the collection inside the Student\_detail database, execute the command as:

```
show collections
```

The command is executed as shown in Figure 2.17.

```
Student_detail> show collections
Studentinfo
Student_detail> |
```

**Figure 2.17: Show Collections**

### 2.4.3 Insert Document

Documents can be inserted into a collection in a database. MongoDB provides options to insert a single document or multiple documents into a collection.

#### Insert a Single Document

The syntax to insert a single document into a collection is:

```
db.collection.insertOne()
```

Consider that the user wants to insert a document into the Studentname collection in the Student\_detail database.

To perform this task, execute the commands as:

```
db.createCollection("Studentname")
db.Studentname.insertOne({ "Name": "Richard"})
```

The command executes as shown in Figure 2.18.

```
Student_detail> db.createCollection("Studentname")
{ ok: 1 }
Student_detail> db.Studentname.insertOne({ "Name": "Richard"})
{
  acknowledged: true,
  insertedId: ObjectId("646219f520cf1016275bd893")
}
```

**Figure 2.18: Insert Single Document**

A document is inserted into the Studentname collection.

The insertedID value in the output is the unique identifier generated by MongoDB for each document. This identifier is of the ObjectId datatype.



---

If the user tries to insert a document into a collection that does not exist in the database, then MongoDB automatically creates the collection.

---

### Insert Multiple Documents

The syntax to insert multiple documents in a collection is:

```
db.collection.insertMany()
```

Consider that the user wants to insert three documents into the Studentmarks collection in the Student\_detail database.

To perform this task, execute the command as:

```
db.Studentmarks.insertMany([
  { Name: "Robert", Age: 16, Subject1: 89, Subject2: 78},
  { Name: "Oliver", Age: 17, Subject1: 78, Subject2: 85} ,
  { Name: "Henry", Age: 15, Subject1: 90, Subject2: 93} ])
```

The command executes as shown in Figure 2.19.

```
Student_detail> db.Studentmarks.insertMany([
... { Name: "Robert", Age: 16, Subject1: 89, Subject2: 78},
... { Name: "Oliver", Age: 17, Subject1: 78, Subject2: 85} ,
... { Name: "Henry", Age: 15, Subject1: 90, Subject2: 93} ])
{
  acknowledged: true,
  insertedIds: {
    '0': ObjectId("64608b0820cf1016275bd88c"),
    '1': ObjectId("64608b0820cf1016275bd88d"),
    '2': ObjectId("64608b0820cf1016275bd88e")
  }
}
```

**Figure 2.19: Insert Many Documents**

#### 2.4.4 Query Document

MongoDB provides commands to retrieve the documents from a collection. The syntax to retrieve all the documents in a collection is:

```
db.collection.find()
```

Consider that the user wants to retrieve all the documents in the Studentmarks collection. To perform this task, execute the command as:

```
db.Studentmarks.find()
```

The command executes as shown in Figure 2.20.

```
Student_detail> db.Studentmarks.find()
[  
  {  
    _id: ObjectId("64608b0820cf1016275bd88c"),  
    Name: 'Robert',  
    Age: 16,  
    Subject1: 89,  
    Subject2: 78  
  },  
  {  
    _id: ObjectId("64608b0820cf1016275bd88d"),  
    Name: 'Oliver',  
    Age: 17,  
    Subject1: 78,  
    Subject2: 85  
  },  
  {  
    _id: ObjectId("64608b0820cf1016275bd88e"),  
    Name: 'Henry',  
    Age: 15,  
    Subject1: 90,  
    Subject2: 93  
  }  
]
```

**Figure 2.20: Retrieve Documents in a Collection**

All the documents in the Studentmarks collection are retrieved and displayed.

#### 2.4.5 Update Document

MongoDB provides two different methods to update documents in a collection. These methods can be used to:

- Update single document
- Update multiple documents

##### Update Single Document

The syntax to update a single document is:

```
db.collection.updateOne(filter, update, options)
```

Table 2.4 describes the parameters of the `updateOne` method.

Parameter	Type	Description
<code>filter</code>	Document	This parameter specifies the selection criteria based on which the update must be performed. If multiple documents satisfy the criteria, then only the first document is updated. This parameter is mandatory.
<code>update</code>	Document	This parameter specifies the change that must be applied to the filtered documents. This parameter is mandatory.
<code>options</code>	Document	This parameter specifies some additional options— <code>upsert</code> , <code>writeconcern</code> , <code>collation</code> , <code>arrayfilters</code> , and <code>hint</code> —for the update method. These options will be explored in detail in later sessions. This parameter is optional.

**Table 2.4: Parameters of `updateOne()` Method**

Consider that the user wants to update the value of `name` field as "David" where the value in the `Name` field is "Robert". To perform this task, execute the command as:

```
db.Studentmarks.updateOne({ "Name": "Robert" },
{ $set: { "Name": "David" } })
```

The command executes as shown in Figure 2.21.

```
Student_detail> db.Studentmarks.updateOne({ "Name": "Robert" }, { $set: { "Name": "David" } })
{
  acknowledged: true,
  insertedId: null,
  matchedCount: 1,
  modifiedCount: 1,
  upsertedCount: 0
}
```

**Figure 2.21: Method to Update Single Document**

The document is updated. To view the updated document, execute the command as:

```
db.Studentmarks.find()
```

The command executes as shown in Figure 2.22.

```
Student_detail> db.Studentmarks.find()
[
  {
    _id: ObjectId("64608b0820cf1016275bd88c"),
    Name: 'David',
    Age: 16,
    Subject1: 89,
    Subject2: 78
  },
  {
    _id: ObjectId("64608b0820cf1016275bd88d"),
    Name: 'Oliver',
    Age: 17,
    Subject1: 78,
    Subject2: 85
  },
  {
    _id: ObjectId("64608b0820cf1016275bd88e"),
    Name: 'Henry',
    Age: 15,
    Subject1: 90,
    Subject2: 93
  }
]
```

**Figure 2.22: Updated Document**

### Update Multiple Documents

The syntax to update multiple documents is:

```
db.collection.updateMany(filter, update, options)
```

Consider that the user wants to update and set the value of Subject2 as 99 in all documents where the value of Subject1 is greater than 80. To perform this task, execute the command as:

```
db.Studentmarks.updateMany({ "Subject1": { $gt: 80 } },
{ $set: { "Subject2": 99 } })
```

The command executes as shown in Figure 2.23.

```
Student_detail> db.Studentmarks.updateMany({ "Subject1":{$gt: 80}}, { $set:{ "Subject2":99}})  
{  
  acknowledged: true,  
  insertedId: null,  
  matchedCount: 2,  
  modifiedCount: 2,  
  upsertedCount: 0  
}
```

**Figure 2.23: Method to Update Many Documents**

The document is updated. In this command, \$gt is a comparison operator that is used to filter out only those documents with Subject1 greater than 80.

To view the updated documents, execute the command as:

```
db.Studentmarks.find()
```

The command executes as shown in Figure 2.24.

```
Student_detail> db.Studentmarks.find()  
[  
  {  
    _id: ObjectId("64617a0420cf1016275bd88f"),  
    Name: 'Robert',  
    Age: 16,  
    Subject1: 89,  
    Subject2: 99  
  },  
  {  
    _id: ObjectId("64617a0420cf1016275bd890"),  
    Name: 'Oliver',  
    Age: 17,  
    Subject1: 78,  
    Subject2: 85  
  },  
  {  
    _id: ObjectId("64617a0420cf1016275bd891"),  
    Name: 'Henry',  
    Age: 15,  
    Subject1: 90,  
    Subject2: 99  
  }  
]
```

**Figure 2.24: Updated Documents**

Here, two documents are updated as they satisfy the specified criteria.

#### 2.4.6 Delete Document

MongoDB provides two different methods to delete documents in a collection. These methods can be used to:

- Delete single document
- Delete many documents

##### Delete Single Document

The syntax to delete single document is:

```
db.collection.deleteOne(filter, options)
```

Table 2.5 describes the parameters of the `deleteOne` method.

Parameter	Type	Description
filter	Document	This parameter specifies the selection criteria based on which the delete must be performed. If an empty criterion is mentioned, then the first document returned in the collection is deleted. This parameter is mandatory.
options	Document	This parameter specifies some additional options—writeconcern, collation, and hint—for the delete method. These options will be explored in detail in later sessions. This parameter is optional.

**Table 2.5: Parameters of `deleteOne` Method**

Consider that the user wants to delete the document where the value of the Name field is "Oliver". To perform this task, execute the command as:

```
db.Studentmarks.deleteOne({ "Name": "Oliver" })
```

The command executes as shown in Figure 2.25.

```
Student_detail> db.Studentmarks.deleteOne({ "Name": "Oliver" })
{ acknowledged: true, deletedCount: 1 }
```

**Figure 2.25: Method to Delete Single Document**

The document where the value of the name field is "Oliver" is deleted. To verify this, execute the command as:

```
db.Studentmarks.find()
```

The command executes as shown in Figure 2.26.

```
student_detail> db.Studentmarks.find()
[
  {
    _id: ObjectId("64617a0420cf1016275bd88f"),
    Name: 'Robert',
    Age: 16,
    Subject1: 89,
    Subject2: 99
  },
  {
    _id: ObjectId("64617a0420cf1016275bd891"),
    Name: 'Henry',
    Age: 15,
    Subject1: 90,
    Subject2: 99
  }
]
```

**Figure 2.26: Output of deleteOne Method**

## Delete Multiple Documents

The syntax to delete multiple documents is:

```
db.collection.deleteMany(filter, options)
```

Table 2.6 describes the parameters of the `deleteMany` method.

Parameter	Type	Description
filter	Document	This parameter specifies the selection criteria based on which the delete must be performed. If an empty criterion is mentioned, then all the documents in the collection are deleted. This parameter is mandatory.
options	Document	This parameter specifies some additional options—writeconcern, collation, and hint—for the delete method. These options will be explored in detail in later sessions. This parameter is optional.

**Table 2.6: Parameters of deleteMany Method**

Consider that the user wants to delete all the documents where the value in the `Age` field is greater than 14. To perform this task, execute the command as:

```
db.Studentmarks.deleteMany({ "Age": { $gt:14 } })
```

The command executes as shown in Figure 2.27.

```
Student_detail> db.Studentmarks.deleteMany({ "Age": { $gt:14 } })
{ acknowledged: true, deletedCount: 2 }
```

**Figure 2.27: Method to Delete Many Documents**

Those documents where the value of the `Age` field is greater than 14 are deleted. To view the output, execute the command as:

```
db.Studentmarks.find()
```

The command executes as shown in Figure 2.28.

```
Student_detail> db.Studentmarks.find()
Student_detail> |
```

**Figure 2.28: Output of `deleteMany` Method**

#### 2.4.7 Drop Collection

MongoDB provides a method to remove a collection from the database. The syntax to drop a collection is:

```
db.collection.drop()
```

Consider that the user wants to drop the collection named `Studentmarks` from the `Student_detail` database. To perform this task, execute the command as:

```
db.Studentmarks.drop()
```

The command executes as shown in Figure 2.29.

```
Student_detail> db.Studentmarks.drop()  
true
```

**Figure 2.29: Dropping a Collection**

The Studentmarks collection is dropped. In MongoDB, when a collection is dropped, all the indexes associated with it are also dropped.

#### 2.4.8 Drop Database

MongoDB provides a method to drop a database. The syntax to drop a database is:

```
db.dropDatabase()
```

This method drops the current database. Consider that the user wants to drop the Student\_detail database. To perform this task, execute the command as:

```
db.dropDatabase()
```

The command is executed as shown in Figure 2.30.

```
Student_detail> db.dropDatabase()  
{ ok: 1, dropped: 'Student_detail' }
```

**Figure 2.30: Method to Drop a Database**

The Student\_detail database is dropped. When a database is dropped, all the collections and data files associated with it are also dropped. Here, the Studentinfo and Studentname collections are also dropped.

## **2.5 Summary**

---

- MongoDB provides a set of tools that helps users to manage the databases in the MongoDB environment.
- To start using the commands in MongoDB, install MongoDB Database Tools and set the environment variable for it.
- A sample dataset can be loaded into MongoDB server using the mongoimport tool.
- MongoDB stores data in BSON format. It supports various datatypes such as string, integer, double, Boolean, null, array, object, objectID, and date.
- MongoDB provides different commands and methods to create and delete databases and collections. It also provides methods to insert, query, update, and delete documents and collections.

## **Test Your Knowledge**

---

1. Which of the following file formats can be imported into MongoDB instance using the mongoimport database tool?
  - a. CSV
  - b. JSON
  - c. DBF
  - d. TSV
  
2. Which of the following commands is used to view the databases on the MongoDB server?
  - a. show databases
  - b. view dbs
  - c. show dbs
  - d. view databases
  
3. Consider that the user wants to create a database named Employee\_detail. Which of the following commands must be used to create the database?
  - a. use Employee\_detail
  - b. create Employee\_detail
  - c. create database Employee\_detail
  - d. use database Employee\_detail
  
4. Which of the following datatypes is used to store embedded documents in MongoDB?
  - a. Array
  - b. Object
  - c. String object
  - d. Array Object
  
5. Which of the following is the correct syntax to delete a collection from the current database?
  - a. db.collection.drop
  - b. db.drop
  - c. db.collection.delete
  - d. db.delete

## **Answers to Test Your Knowledge**

---

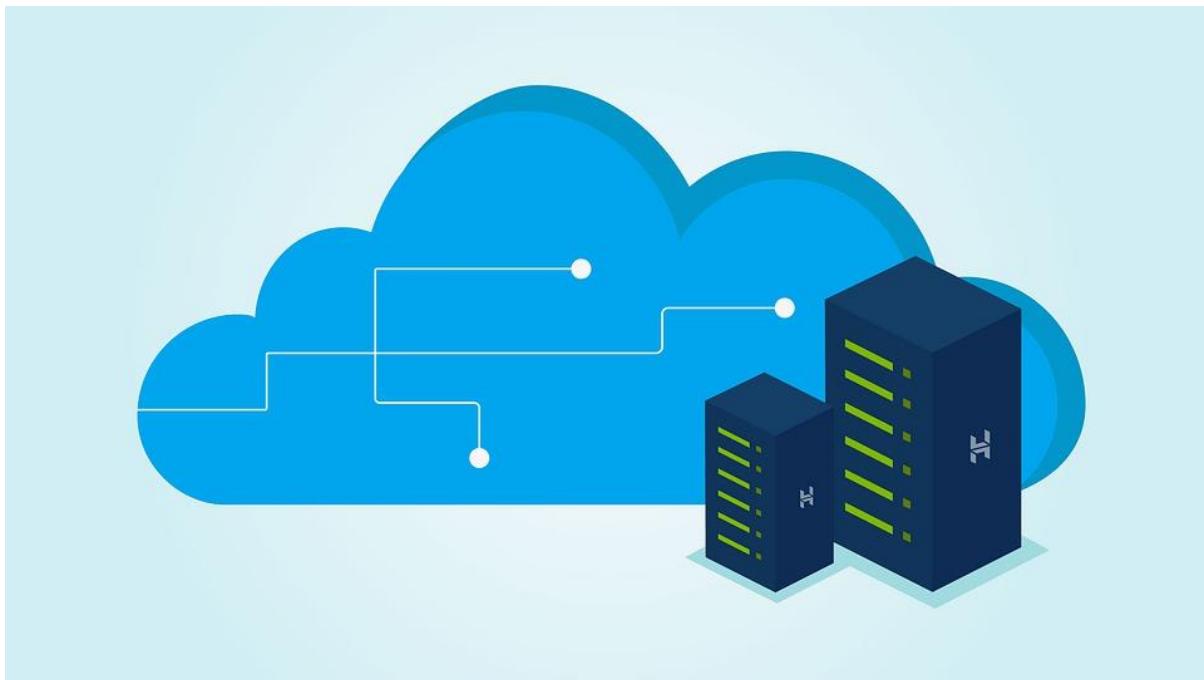
1	a, b, d
2	c
3	a
4	b
5	a

## **Try it Yourself**

---

1. Install MongoDB Database Tools and set the environment variable for MongoDB Database Tools.
2. Load the sample dataset `sample_training` into the MongoDB server.
3. Create a database named `Employee` and a collection named `Employee_detail`.
4. Insert the following documents into the `Employee_detail` collection.

```
[ {  
    Emp_name: "Oliver Smith",  
    Age: 23,  
    Designation: "Manager"  
},  
{  
    Emp_name: "David Michael",  
    Age: 32,  
    Designation: "Software Engineer"  
}]
```
5. Update the designation of "Oliver Smith" to "Accountant".
6. Delete the document in `Employee_detail` collection where the name of the employee is "David Michael".
7. Drop the collection `Employee_detail`.
8. Drop the database `Employee`.



## SESSION 3

### MONGODB OPERATORS

---

#### Learning Objectives

In this session, students will learn to:

- Describe the types of MongoDB operators
- Explain how to use the query and projection operators
- Explain how to modify field and array data using update operators

MongoDB provides various operators which enable users to effectively interact with the database.

This session describes the types of MongoDB operators and their use in general. This session also describes in detail the function of the more commonly used operators with real-life examples.

#### **3.1 Introduction to MongoDB Operators**

---

Operators are reserved words or symbols that instruct the compiler or interpreter to perform specific mathematical or logical operations on the dataset as required.

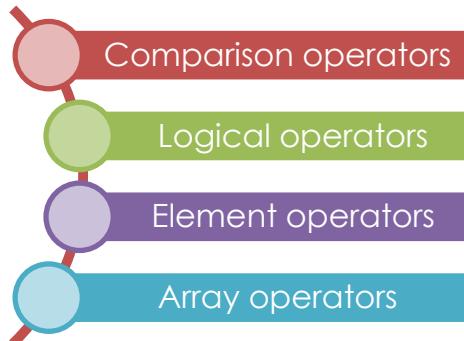
Some of the types of operators in MongoDB are:

- Query operators
- Projection operators
- Update operators
- Miscellaneous operators

## 3.2 Query Operators

---

Query operators are used to retrieve data from a database. Different types of query operators are:



### 3.2.1 Comparison Operators

Comparison operators compare the field's value with the specified value and return the documents that match the value. Table 3.1 describes the comparison operators in MongoDB.

Operator	Description
\$eq	Used to retrieve, update, or delete documents where the field value is equal to a specified value
\$gt	Used to retrieve, update, or delete documents where the field value is greater than a specified value
\$gte	Used to retrieve, update, or delete documents where the field value is greater than or equal to a specified value
\$in	Used to retrieve, update, or delete documents where the field value is any of the values specified in an array
\$lt	Used to retrieve, update, or delete documents where the field value is less than a specified value
\$lte	Used to retrieve, update, or delete documents where the field value is less than or equal to a specified value
\$ne	Used to retrieve, update, or delete documents where the field value is not equal to a specified value
\$nin	Used to retrieve, update, or delete documents where the field value is none of the values specified in an array.

**Table 3.1: Comparison Operators in MongoDB**

Let us take a look at how to use some of these comparison operators.

### \$eq Operator

This operator checks whether the value of a field equals the specified value. The syntax for this operator is:

```
{ <field>: { $eq: <value> } }
```

Let us take a look at how to use this operator.

Figure 3.1 shows the documents stored in the `accounts` collection in the sample database named `sample_analytics`.

```
sample_analytics> db.accounts.find()
[
  {
    _id: ObjectId("5ca4bbc7a2dd94ee5816238c"),
    account_id: 371138,
    limit: 9000,
    products: [ 'Derivatives', 'InvestmentStock' ]
  },
  {
    _id: ObjectId("5ca4bbc7a2dd94ee5816238d"),
    account_id: 557378,
    limit: 10000,
    products: [ 'InvestmentStock', 'Commodity', 'Brokerage', 'CurrencyService' ]
  },
  {
    _id: ObjectId("5ca4bbc7a2dd94ee5816238e"),
    account_id: 198100,
    limit: 10000,
    products: [ 'Derivatives', 'CurrencyService', 'InvestmentStock' ]
  },
  {
    _id: ObjectId("5ca4bbc7a2dd94ee58162390"),
    account_id: 278603,
    limit: 10000,
    products: [ 'Commodity', 'InvestmentStock' ]
  }
]
```

**Figure 3.1: Documents in the Accounts Collection**

Now, consider that the user wants to view the `accounts` which have the `limit` value as 9000. To do so, the user can execute a query that uses the `$eq` operator as:

```
db.accounts.find({limit:{$eq:9000}})
```

Figure 3.2 shows the output of this query.

```
sample_analytics> db.accounts.find({limit:{$eq:9000}})
[{"_id": ObjectId("5ca4bbc7a2dd94ee5816238c"), "account_id": 371138, "limit": 9000, "products": ["Derivatives", "InvestmentStock"]}, {"_id": ObjectId("5ca4bbc7a2dd94ee58162392"), "account_id": 794875, "limit": 9000, "products": ["InvestmentFund", "InvestmentStock"]}, {"_id": ObjectId("5ca4bbc7a2dd94ee581623e0"), "account_id": 161714, "limit": 9000, "products": ["Commodity", "CurrencyService", "Derivatives", "InvestmentFund", "InvestmentStock"]}, {"_id": ObjectId("5ca4bbc7a2dd94ee581623ee")}]
```

**Figure 3.2: Output of the Query with \$eq Operator**

### \$lt Operator

This operator checks whether the value of a field is less than the specified value. The syntax for this operator is:

```
{ <field>: { $lt: <value> } }
```

For example, consider that in the `accounts` collection, the user wants to view the documents where the `limit` value is less than 5000. To do so, the user can execute a query that uses the `$lt` operator as:

```
db.accounts.find({limit:{$lt:5000}})
```

Figure 3.3 shows the output of the query.

```
sample_analytics> db.accounts.find({limit:{$lt:5000}})  
[  
  {  
    _id: ObjectId("5ca4bbc7a2dd94ee58162661"),  
    account_id: 417993,  
    limit: 3000,  
    products: [ 'InvestmentStock', 'InvestmentFund' ]  
  },  
  {  
    _id: ObjectId("5ca4bbc7a2dd94ee581626ad"),  
    account_id: 113123,  
    limit: 3000,  
    products: [ 'CurrencyService', 'InvestmentStock' ]  
  }  
]  
sample_analytics>
```

**Figure 3.3: Output of the Query with \$lt Operator**

### \$in Operator

This operator checks whether the value of a field matches any of the values in the specified array. The syntax for this operator is:

```
{} field: { $in: [<value1>, <value2>, ... <valueN> ] } }
```

For example, consider that from the `accounts` collection, the user wants to view only the documents where the product is either a 'Commodity' or an 'InvestmentStock'. To do so, the user can execute a query that uses the `$in` operator as:

```
db.accounts.find({products:  
{$in: ['Commodity', 'InvestmentStock']} })
```

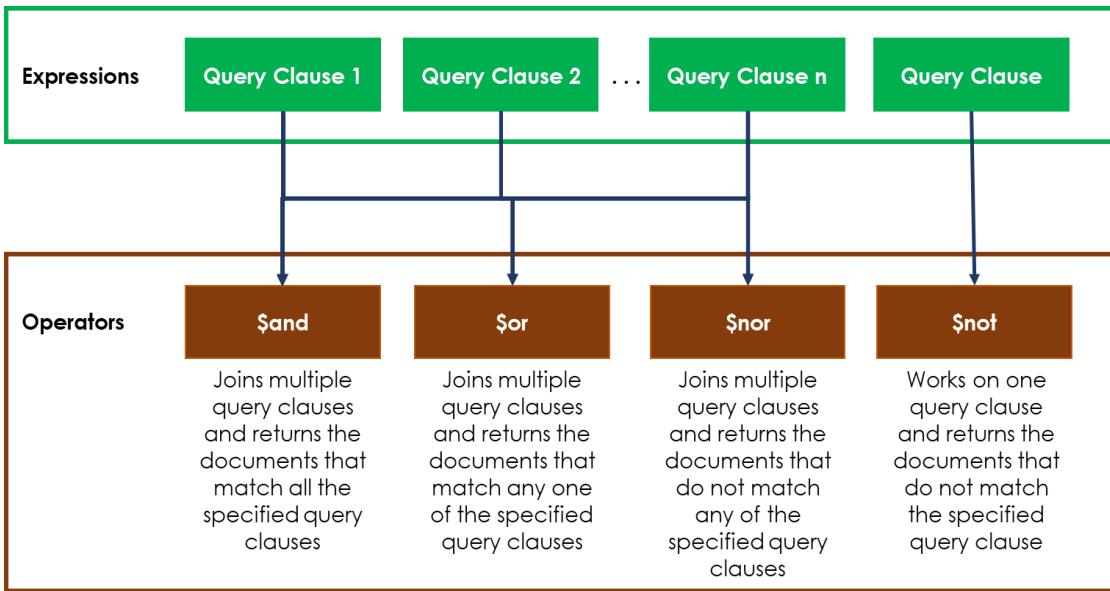
Figure 3.4 shows the output of the query.

```
'sample_analytics> db.accounts.find({products: {$in:[ 'Commodity', 'InvestmentStock']}})'
[  
  {  
    _id: ObjectId("5ca4bbc7a2dd94ee5816238c"),  
    account_id: 371138,  
    limit: 9000,  
    products: [ 'Derivatives', 'InvestmentStock' ]  
  },  
  {  
    _id: ObjectId("5ca4bbc7a2dd94ee5816238d"),  
    account_id: 557378,  
    limit: 10000,  
    products: [ 'InvestmentStock', 'Commodity', 'Brokerage', 'CurrencyService' ]  
  },  
  {  
    _id: ObjectId("5ca4bbc7a2dd94ee5816238e"),  
    account_id: 198100,  
    limit: 10000,  
    products: [ 'Derivatives', 'CurrencyService', 'InvestmentStock' ]  
  },  
  {  
    _id: ObjectId("5ca4bbc7a2dd94ee58162390"),  
    account_id: 278603,  
    limit: 10000,  
    products: [ 'Commodity', 'InvestmentStock' ]  
  },  
  {  
    _id: ObjectId("5ca4bbc7a2dd94ee58162392"),  
    account_id: 198100,  
    limit: 10000,  
    products: [ 'Derivatives', 'CurrencyService', 'InvestmentStock' ]  
  }]
```

**Figure 3.4: Output of the Query with \$in Operator**

### 3.2.2 Logical Operators

Logical operators evaluate multiple expressions and return a Boolean value (true, false). Four logical operators are: \$and, \$or, \$nor, and \$not:



## \$and Operator

This operator, as the name suggests, performs a logical AND operation. This means that the operator lists only the documents that match all the expressions in the specified array.

The syntax for this operator is:

```
{ $and: [ { <expression1> }, { <expression2> } , ... ,  
      { <expressionN> } ] }
```

Consider that from the accounts collection, the user wants to view only the documents where the limit is greater than or equal to 5000 and less than 7000. To do so, the user can execute a query that uses the \$and operator as:

```
db.accounts.find({$and: [{limit:{$gte:5000}},  
{limit:{$lt:7000}}] })
```

Figure 3.5 shows the output of this query.

```
sample_analytics> db.accounts.find({$and: [{limit:{$gte:5000}}, {limit:{$lt:7000}}]})
[
  {
    _id: ObjectId("5ca4bbc7a2dd94ee5816272e"),
    account_id: 170980,
    limit: 5000,
    products: [
      'InvestmentFund',
      'Brokerage',
      'CurrencyService',
      'Derivatives',
      'InvestmentStock'
    ]
  }
]
```

**Figure 3.5: Output of the Query with \$and Operator**

### \$or Operator

The \$or operator performs a logical OR operation. This operator lists documents that satisfy at least one of the expressions in the specified array.

The syntax for this operator is:

```
{ $or: [ { <expression1> }, { <expression2> }, ... , {
  <expressionN> } ] }
```

Consider that from the account collection, the user wants to view all the documents where the limit is less than or equal to 5000 or equal to 7000. To do, the user can execute a query that uses the \$or operator as:

```
db.accounts.find({$or: [{limit:{$lte:5000}}, {limit:{$eq:7000}}] })
```

Figure 3.6 shows the output of the query.

```
sample_analytics> db.accounts.find({$or: [{limit:{$lte:5000}}, {limit:{$eq:7000}}]})  
[  
  {  
    _id: ObjectId("5ca4bbc7a2dd94ee58162458"),  
    account_id: 852986,  
    limit: 7000,  
    products: [  
      'Derivatives',  
      'Commodity',  
      'CurrencyService',  
      'InvestmentFund',  
      'InvestmentStock'  
    ]  
  },  
  {  
    _id: ObjectId("5ca4bbc7a2dd94ee5816247a"),  
    account_id: 777752,  
    limit: 7000,  
    products: [  
      'CurrencyService',  
      'Brokerage',  
      'Commodity',  
      'InvestmentFund',  
      'InvestmentStock'  
    ]  
  },  
  {
```

**Figure 3.6: Output of the Query with \$or Operator**

### \$nor Operator

This operator performs a logical NOR operation. It lists the documents that do not match any of the expressions in the specified array of expressions.

The syntax for this operator is:

```
{ $nor: [ { <expression1> }, { <expression2> }, ... {  
  <expressionN> } ] }
```

For example, consider that from the `accounts` collection, the user wants to view all the documents where the `limit` is neither 9000 nor 10000. To do so, the users can execute a query that uses the `$nor` operator as:

```
db.accounts.find({$nor: [{limit:9000}, {limit:10000}]})
```

Figure 3.7 shows the output of this query.

```
sample_analytics> db.accounts.find({$nor: [{limit:9000}, {limit:10000}]}  
[  
  {  
    _id: ObjectId("5ca4bbc7a2dd94ee58162458"),  
    account_id: 852986,  
    limit: 7000,  
    products: [  
      'Derivatives',  
      'Commodity',  
      'CurrencyService',  
      'InvestmentFund',  
      'InvestmentStock'  
    ]  
  },  
  {  
    _id: ObjectId("5ca4bbc7a2dd94ee5816247a"),  
    account_id: 777752,  
    limit: 7000,  
    products: [  
      'CurrencyService',  
      'Brokerage',  
      'Commodity',  
      'InvestmentFund',  
      'InvestmentStock'  
    ]  
  },  
]
```

**Figure 3.7: Output of the Query with \$nor Operator**

### \$not Operator

This operator performs a logical NOT operation. It lists the documents that do not match the specified expression.

The syntax for this operator is:

```
{ field: { $not: { <operator-expression> } } }
```

For example, consider that from the accounts collection, the user wants to view all the documents where limit is not greater than or equal to 5000.

To do so, the user can execute a query that uses the \$not operator as:

```
db.accounts.find( {limit: {$not:{$gte:5000}}})
```

Figure 3.8 shows the output of this query.

```
sample_analytics> db.accounts.find( {limit: {$not:{$gte:5000}}})
[
  {
    _id: ObjectId("5ca4bbc7a2dd94ee58162661"),
    account_id: 417993,
    limit: 3000,
    products: [ 'InvestmentStock', 'InvestmentFund' ]
  },
  {
    _id: ObjectId("5ca4bbc7a2dd94ee581626ad"),
    account_id: 113123,
    limit: 3000,
    products: [ 'CurrencyService', 'InvestmentStock' ]
  }
]
sample_analytics>
```

**Figure 3.8: Output of the Query with \$not Operator**

### 3.2.3 Element Operators

Element operators are used to check if a particular field is available in a MongoDB document or if a particular field with a specified type is available in the document. These operators are used with either of the two values: true or false. If these operators are used with the true value, the query will return all the documents that include the specified field or the specified field with the specified datatype. If used with the false value, the query will return all the documents that do not include the specified field or the specified field with the specified datatype.

MongoDB supports two element operators:

- \$exists
- \$type

#### **\$exists Operator**

The \$exists operator checks if the specified field is present in a document. If it is used with the true value, \$exists returns the documents that contain the specified field, including documents where the field value is null. If \$exists is used with the false value, \$exists returns only the documents that do not contain the specified field.

The syntax for this operator is:

```
{ field: { $exists: <boolean> } }
```

In this syntax, <Boolean> takes two values: true or false.

Before understanding how this operator works, let us create a database named `Product_detail` and a collection named `product_delivery`. Let us also insert the five documents into the `product_delivery` collection using the query as:

```
db.product_delivery.insertMany(  
  [ { "_id" : 1, address : "2130 Messy Way", zipCode :  
    "90698345", delivered:["Documents","Food","electronics",  
    "Books"],Feedback:[{"product":"xyz","score":6}, {"product":"abc"  
    ,"score":7} ] ,  
    { "_id" : 2, address: "156 Lunar way garden",  
delivered:["Food","electronics"],Feedback:[ {"product":"xyz", "s  
core":5}, {"product":"abc", "score":4} ] ,  
      { "_id" : 3, address : "456 Penny way Palace", zipCode:  
3921412,  
delivered:["electronics"],Feedback:[ {"product":"xyz", "score":8  
}, {"product":"abc", "score":9} ] ,  
      { "_id" : 4, address : "155 Solar Ring  
road",delivered:["Books"],Feedback:[ {"product":"xyz", "score":5  
}, {"product":"abc", "score":2} ] ,  
      { "_id" : 5, address : "134 Julie Garden", zipCode :  
["834847278",  
"1893289032"],delivered:["Food","electronics"],Feedback:[ {"pro  
duct":"xyz", "score":4}, {"product":"abc", "score":8} ] }  
  ] )
```

Figure 3.9 shows the output of this query.

```
Product_detail> db.product_delivery.insertMany(  
... [  
...   { "_id" : 1, address : "2130 Messy Way", zipCode : "90698345", delivered:["Documents","Food","electronics", "books"],Feedback:[{"product":"xyz","score":6}, {"product":"abc","score":7}] },  
...   { "_id" : 2, address: "156 Lunar way garden", delivered:["Food", "electronics"],Feedback:[{"product":"xyz","score":5}, {"product":"abc","score":4}] },  
...   { "_id" : 3, address : "456 Penny way Palace", zipCode: 3921412, delivered:["electronics"],Feedback:[{"product":"xyz","score":8}, {"product":"abc","score":9}] },  
...   { "_id" : 4, address : "155 Solar Ring road",delivered:["Books"],Feedback:[{"product":"xyz","score":5}, {"product":"abc","score":2}] },  
...   { "_id" : 5, address : "134 Julie Garden", zipCode : ["834847278", "1893289032"],delivered:["Food", "electronics"],Feedback:[{"product":"xyz","score":4}, {"product":"abc","score":8}] }  
... ]  
... )  
{  
  acknowledged: true,  
  insertedIds: { '0': 1, '1': 2, '2': 3, '3': 4, '4': 5 }  
}
```

**Figure 3.9: Output of the insertMany Query**

Now, consider that from the product\_delivery collection, the user wants to view documents that do not have a zip code. To do, so, the user can execute a query that uses the \$exists operator as:

```
db.product_delivery.find({"zipCode":{$exists:false}})
```

Figure 3.10 shows the output of this query.

```
Product_detail> db.product_delivery.find({"zipCode":{$exists:false}})  
[  
  {  
    _id: 2,  
    address: '156 Lunar way garden',  
    delivered: [ 'Food', 'electronics' ],  
    Feedback: [ { product: 'xyz', score: 5 }, { product: 'abc', score: 4 } ]  
  },  
  {  
    _id: 4,  
    address: '155 Solar Ring road',  
    delivered: [ 'Books' ],  
    Feedback: [ { product: 'xyz', score: 5 }, { product: 'abc', score: 2 } ]  
  }  
]
```

**Figure 3.10: Output of the Query with \$exists Operator**

## \$type Operator

This operator returns the documents where the datatype of the specified field in the document matches the specified datatype. The \$type operator is quite useful when dealing with highly unstructured data where the collection is complex and the prediction of the datatypes of any field is difficult.

The syntax for this operator is:

```
{ field: { $type: <BSON type> } }
```

The \$type operator accepts aliases and numbers in place of the datatype. Table 3.2 lists the datatype and its corresponding number and alias.

Datatype	Number	Alias
Double	1	"double"
String	2	"string"
Object	3	"object"
Array	4	"array"
Binary data	5	"binData"
ObjectId	7	"objectId"
Boolean	8	"bool"
Date	9	"date"
Null	10	"null"
Regular Expression	11	"regex"
JavaScript	13	"javascript"
32-bit integer	16	"int"
Timestamp	17	"timestamp"
64-bit integer	18	"long"
Decimal128	19	"decimal"
Min key	-1	"minKey"
Max key	127	"maxKey"

**Table 3.2: Datatypes and their Corresponding Numbers and Aliases**

Now, consider that from the product\_delivery collection, the user wants to view the zipCode fields which have the string type. To do so, the user can execute a query that uses the \$type operator as:

```
db.product_delivery.find( { "zipCode" : { $type : 2}})
```

Figure 3.11 shows the output of the query.

```
Product_detail> db.product_delivery.find( { "zipCode" : { $type : 2}})
[
  {
    _id: 1,
    address: '2130 Messy Way',
    zipCode: '90698345',
    delivered: [ 'Documents', 'Food', 'electronics', 'Books' ],
    Feedback: [ { product: 'xyz', score: 6 }, { product: 'abc', score: 7 } ]
  },
  {
    _id: 5,
    address: '134 Julie Garden',
    zipCode: [ '834847278', '1893289032' ],
    delivered: [ 'Food', 'electronics' ],
    Feedback: [ { product: 'xyz', score: 4 }, { product: 'abc', score: 8 } ]
  }
]
```

**Figure 3.11: Output of the Query with \$type Operator**

Now, consider that from the `product_delivery` collection, the user wants to view the `zipCode` fields which are of the string type. To do so, the user can execute a query that uses the `$type` operator as:

```
db.product_delivery.find( { "zipCode" : { $type : 16}})
```

Figure 3.12 shows the output of the query.

```
Product_detail> db.product_delivery.find( { "zipCode" : { $type : 16}})
[
  {
    _id: 3,
    address: '456 Penny way Palace',
    zipCode: 3921412,
    delivered: [ 'electronics' ],
    Feedback: [ { product: 'xyz', score: 8 }, { product: 'abc', score: 9 } ]
  }
]
```

**Figure 3.12: Output of the Query with \$type Operator**

### 3.2.4 Array Operators

Array operators are used when the user wants to filter data by specifying conditions for array fields.

There are three array operators:

#### \$all

Used to retrieve, update, or delete documents that have an array field that has all the elements specified in the query

#### \$elemMatch

Used to retrieve, update, or delete documents that have an array field that has at least one of the elements specified in the query

#### \$size

Used to retrieve, update, or delete documents that have an array field with the same size as the size specified in the query

#### \$all Operator

The syntax for this operator is:

```
{ <field>: { $all: [ <value1> , <value2> ... ] } }
```

For example, consider that from the product\_delivery collection, the user wants to view the documents that include the value 'Food' and 'electronics' in the delivered field. To do so, the user can execute a query that uses the \$all operator as:

```
db.product_delivery.find({delivered:  
{$all:['Food', 'electronics']} })
```

Figure 3.13 shows the output of the query.

```
Product_detail> db.product_delivery.find({delivered: {$all:['Food', 'electronics']}})
[
  {
    _id: 1,
    address: '2130 Messy Way',
    zipCode: '90698345',
    delivered: [ 'Documents', 'Food', 'electronics', 'Books' ],
    Feedback: [ { product: 'xyz', score: 6 }, { product: 'abc', score: 7 } ]
  },
  {
    _id: 2,
    address: '156 Lunar way garden',
    delivered: [ 'Food', 'electronics' ],
    Feedback: [ { product: 'xyz', score: 5 }, { product: 'abc', score: 4 } ]
  },
  {
    _id: 5,
    address: '134 Julie Garden',
    zipCode: [ '834847278', '1893289032' ],
    delivered: [ 'Food', 'electronics' ],
    Feedback: [ { product: 'xyz', score: 4 }, { product: 'abc', score: 8 } ]
  }
]
```

**Figure 3.13: Output of the Query with \$all Operator**

### \$elemMatch Operator

The syntax for this operator is:

```
{ <field>: { $elemMatch: { <query1>, <query2>, ... } } }
```

For example, consider that from the `product_delivery` collection, the user wants to view the documents that have the value `product: "xyz"` or `score` greater than `8` in the `Feedback` field. To do so, the user can execute a query that uses the `$elemMatch` operator as:

```
db.product_delivery.find( { Feedback: { $elemMatch:
{ product: "xyz", score: { $gte: 8} } } })
```

Figure 3.14 shows the output of the query.

```
Product_detail> db.product_delivery.find( { Feedback: { $elemMatch: { product: "xyz", score: { $gte: 8} } } })  
[  
  {  
    _id: 3,  
    address: '456 Penny way Palace',  
    zipCode: 3921412,  
    delivered: [ 'electronics' ],  
    Feedback: [ { product: 'xyz', score: 8 }, { product: 'abc', score: 9 } ]  
  }  
]
```

**Figure 3.14: Output of the Query with \$elemMatch Operator**

### \$size Operator

The syntax for this operator is:

```
{ <field>: { $size: n} }
```

For example, consider that from the `product_delivery` collection, the user wants to view the documents that have four products in the `delivered` array. To do so, the user can execute a query that uses the `$size` operator as:

```
db.product_delivery.find({delivered:{$size:4}})
```

Figure 3.15 shows the output of the query.

```
Product_detail> db.product_delivery.find({delivered:{$size:4}})  
[  
  {  
    _id: 1,  
    address: '2130 Messy Way',  
    zipCode: '90698345',  
    delivered: [ 'Documents', 'Food', 'electronics', 'Books' ],  
    Feedback: [ { product: 'xyz', score: 6 }, { product: 'abc', score: 7 } ]  
  }  
]
```

**Figure 3.15: Output of the Query with \$size Operator**

## 3.3 Projection Operators

The `find()` operator in MongoDB returns all the data in the collection that matches the specified condition. What if the user wants to view specific fields of data from the collection? This is where project operators come in handy.

Projection operators can be used to specify the fields or elements to display in the query result based on the condition specified in the query.

Four projection operators in MongoDB are:

\$	Used to return the first element in an array that matches the condition specified in the query or the first element in the array, if no query condition is specified
\$elemMatch	Used to return the first element in an array that matches the condition specified by the \$elemMatch operator
\$meta	Used to retrieve the metadata related to a document in a collection
\$slice	Used to specify the number of elements that must be returned by the query

## \$ Operator

The syntax for this operator is:

```
db.collection.find( { <array>: <condition> ... },  
                    { "<array>.$": 1 } )
```

To understand how this operator can be used, let us create a database named `student_detail` and switch to that database as shown in Figure 3.16.

```
sample_analytics> use Student_detail  
switched to db Student_detail
```

**Figure 3.16: Switching to the `Student_detail` Database**

Now, let us create a collection named `Studentmarks` in the database. Next, let us insert documents into the `Studentmarks` collection using the query as:

```
db.Studentmarks.insertMany([
    { name: "Robert", age: 16, sub:1, marks:[89,78,90]}, 
    { name: "Oliver", age: 17, sub:1, marks:[78,85,89]}, 
    { name: "Henry", age: 15, sub:1, marks:[88,89,76]}, 
    { name: "David", age: 16, sub:2, marks:[86,84,66]}])
```

Figure 3.17 shows the output of the query.

```
Student_detail> db.Studentmarks.insertMany([
...     { name: "Robert", age: 16, sub:1, marks:[89,78,90]}, 
...     { name: "Oliver", age: 17, sub:1, marks:[78,85,89]}, 
...     { name: "Henry", age: 15, sub:1, marks:[88,89,76]}, 
...     { name: "David", age: 16, sub:2, marks:[86,84,66]}])
{
    acknowledged: true,
    insertedIds: {
        '0': ObjectId("64525fc59a2bc07eed14823d"),
        '1': ObjectId("64525fc59a2bc07eed14823e"),
        '2': ObjectId("64525fc59a2bc07eed14823f"),
        '3': ObjectId("64525fc59a2bc07eed148240")}
```

**Figure 3.17: Output of the `insertMany` Query**

Now consider that the user wants to view the first element that is greater than or equal to 85 for the `marks` field, where the value of `sub` is 1. The user also wants to view the `name` field and hide the `_id` field. To do so, the user can execute a query that uses the `$` projection operator as:

```
db.Studentmarks.find( { sub: 1, marks:
{ $gte: 85 }
}, {_id:0, name:1, "marks.$": 1} )
```




---

In this query, `"marks.$": 1` ensures that the `marks` column is displayed and `_id:0` ensures that the `ID` column is hidden in the query result.

---

Figure 3.18 shows the output of this query.

```
Student_detail> db.Studentmarks.find( { sub: 1, marks: { $gte: 85 } }, { _id:0, name:1, "marks.$": 1 } )
[
  { name: 'Robert', marks: [ 89 ] },
  { name: 'Oliver', marks: [ 85 ] },
  { name: 'Henry', marks: [ 88 ] }
]
```

**Figure 3.18: Output of the Query with \$ Operator**

### \$slice Operator

The syntax for this operator is:

```
db.collection.find( <query>,
{ <arrayField>: { $slice: <number> } } );
```

For example, consider that the user wants to view the first two elements of the marks field from the document where the name field has the value “Robert”. To do so, the user can execute a query that uses the \$slice operator as:

```
db.Studentmarks.find({ "name": "Robert" },
{ "marks": { $slice: 2 }})
```

Figure 3.19 shows the output of this query.

```
Student_detail> db.Studentmarks.find({ "name": "Robert" }, { "marks": { $slice: 2 } })
[
  {
    _id: ObjectId("64525fc59a2bc07eed14823d"),
    name: 'Robert',
    age: 16,
    sub: 1,
    marks: [ 89, 78 ]
  }
]
Student_detail>
```

**Figure 3.19: Output of the Query with \$slice Operator**

## 3.4 Update Operators

Update operators help users modify documents in a collection. There are separate update operators for modifying fields and arrays.

### 3.4.1 Field Update Operators

There are nine operators in MongoDB that allow updating the values of fields in the documents:

#### **\$currentDate**

- Used when the value in a field must be set to the current date

#### **\$inc**

- Used when the values in a field must be incremented by a specific number

#### **\$min**

- Used when the values in a field must be updated only if they are more than a specific value

#### **\$max**

- Used when the values in a field must be updated only if they are less than a specific value

#### **\$mul**

- Used when the values in a field must be multiplied by a specific value

#### **\$rename**

- Used when a field must be renamed

#### **\$set**

- Used when the value of a field in the document must be set

#### **\$unset**

- Used when a specific field must be removed from the document

#### **\$setOnInsert**

- Used when the value of a field must be set only if a new document is inserted as a result of the update operation

#### **\$currentDate Operator**

The syntax for this operator is:

```
{ $currentDate: { <field1>: <typeSpecification1>, ... } }
```

In this syntax, `<typeSpecification>` can be set to `true` if the user wants to input the current date in the `Date` type. If the user wants to insert the current date as a timestamp or a date type, then the user can set `<typeSpecification>` to `timestamp` or `date`, respectively. Note that these values must be used in lower case.



---

If the field specified in the query does not exist in the document, then this operator will add the field to the document.

---

For example, consider that the user wants to set the `Exam_Date` field in the `Studentmarks` collection to the current date as a Date. To do so, the user can execute and update query that uses the `$currentDate` operator as:

```
db.Studentmarks.updateOne({"name": "David"},  
{$currentDate: {Exam_Date: true}})
```

Figure 3.20 shows the output of this query.

```
Student_detail> db.Studentmarks.updateOne({"name": "David"}, {$currentDate: {Exam_Date: true}})  
{  
  acknowledged: true,  
  insertedId: null,  
  matchedCount: 1,  
  modifiedCount: 1,  
  upsertedCount: 0  
}
```

**Figure 3.20: Output of the Query with `$currentDate` Operator**

To view the documents in the updated `Studentmarks` collection, the user can execute the query as:

```
db.Studentmarks.find()
```

Figure 3.21 shows the output of this query.

```
Student_detail> db.Studentmarks.find()
[
  {
    _id: ObjectId("646f510b92732287e845d80d"),
    name: 'Robert',
    age: 16,
    sub: 1,
    marks: [ 89, 78, 90 ]
  },
  {
    _id: ObjectId("646f510b92732287e845d80e"),
    name: 'Oliver',
    age: 17,
    sub: 1,
    marks: [ 78, 85, 89 ]
  },
  {
    _id: ObjectId("646f510b92732287e845d80f"),
    name: 'Henry',
    age: 15,
    sub: 1,
    marks: [ 88, 89, 76 ]
  },
  {
    _id: ObjectId("646f510b92732287e845d810"),
    name: 'David',
    age: 16,
    sub: 2,
    marks: [ 86, 84, 66 ],
    Exam_Date: ISODate("2023-05-25T12:14:19.762Z")
  }
]
```

Figure 3.21: Updated Studentmarks Collection

## \$inc Operator

The syntax for this operator is:

```
{ $inc: { <field1>: <amount1>, <field2>:
          <amount2>, ... } }
```

For example, consider that the user wants to increment the value in the age field by 2 where the name is "Henry". To do so, the user can execute a query that uses the \$inc operator as:

```
db.Studentmarks.updateOne( { name: "Henry" },
{ $inc: { age: 2 } })
```

Figure 3.22 shows the output of this query.

```
Student_details> db.Studentmarks.updateOne( { name: "Henry" }, { $inc: { age: 2 } })
{
  acknowledged: true,
  insertedId: null,
  matchedCount: 1,
  modifiedCount: 1,
  upsertedCount: 0
}
```

**Figure 3.22: Output of the Query with \$inc Operator**

To view the updated document, the user can execute the query as:

```
db.Studentmarks.find({ "name": "Henry" })
```

Figure 3.23 shows the output of this query.

```
Student_details> db.Studentmarks.find({ "name": "Henry" })
[
  {
    _id: ObjectId("646f55e160c2b086a2c39e27"),
    name: 'Henry',
    age: 17,
    sub: 1,
    marks: [ 88, 89, 76 ]
  }
]
```

**Figure 3.23: Updated Document**

### \$set Operator

The syntax for this operator is:

```
{ $set: { <field1>: <value1>, ... } }
```

For example, consider that the user wants to add a new field named `credit` to the `Studentmarks` collection. To do this, the user can execute an update query that uses the `$set` operator as:

```
db.Studentmarks.updateMany( {}, [ { $set: { credit: 0 } } ] )
```

Figure 3.24 shows the output of this query.

```
Student_details> db.Studentmarks.updateMany( {}, [ { $set: { credit:0 } } ] )
{
  acknowledged: true,
  insertedId: null,
  matchedCount: 4,
  modifiedCount: 4,
  upsertedCount: 0
}
```

**Figure 3.24: Output of the Query with \$set Operator**

To view the updated document, the user can execute the query as:

```
db.Studentmarks.find()
```

Figure 3.25 shows the output of this query.

```
Student_details> db.Studentmarks.find()
[
  {
    _id: ObjectId("646f5c0760c2b086a2c39e29"),
    name: 'Robert',
    age: 16,
    sub: 1,
    marks: [ 89, 78, 90 ],
    credit: 0
  },
  {
    _id: ObjectId("646f5c0760c2b086a2c39e2a"),
    name: 'Oliver',
    age: 17,
    sub: 1,
    marks: [ 78, 85, 89 ],
    credit: 0
  },
  {
    _id: ObjectId("646f5c0760c2b086a2c39e2b"),
    name: 'Henry',
    age: 17,
    sub: 1,
    marks: [ 88, 89, 76 ],
    credit: 0
  },
  {
    _id: ObjectId("646f5c0760c2b086a2c39e2c"),
    name: 'David',
    age: 16,
    sub: 2,
    marks: [ 86, 84, 66 ],
    Exam_Date: ISODate("2023-05-25T13:02:05.079Z"),
    credit: 0
  }
]
```

**Figure 3.25: Updated Document**

## \$rename Operator

The syntax for this operator is:

```
{$rename: { <field1>: <newName1>, <field2>:  
          <newName2>, ... } }
```



The new field name must be different from the existing field name.

For example, consider that in the `Studentmarks` collection the user wants to change the name of `credit` field to `unit`. To do this, the user can execute an update query that uses the `$rename` operator as:

```
db.Studentmarks.updateOne( { name: "David" },  
                           { $rename: { 'credit': 'unit' } })
```

Figure 3.26 shows the output of this query.

```
Student_detail> db.Studentmarks.updateOne( { name: "David" }, { $rename: { 'credit': 'unit' } })  
{  
  acknowledged: true,  
  insertedId: null,  
  matchedCount: 1,  
  modifiedCount: 1,  
  upsertedCount: 0  
}
```

**Figure 3.26: Output of the Query with \$rename Operator**

To view the updated document, the user can execute the query as:

```
db.Studentmarks.find( { "name": "David" } )
```

Figure 3.27 shows the output of this query.

```
Student_details> db.Studentmarks.find({"name": "David"})
[{"_id": ObjectId("646f5c0760c2b086a2c39e2c"),
  "name": "David",
  "age": 16,
  "sub": 2,
  "marks": [ 86, 84, 66 ],
  "Exam_Date": ISODate("2023-05-25T13:02:05.079Z"),
  "unit": 0}
```

**Figure 3.27: Updated Document**

### 3.4.2 Array Update Operators

The array update operators supported by MongoDB are:

\$	• Used to update the first element in the specified array that matches the condition specified in the query
\$[]	• Used to update all the elements in the specified array for all the documents that match the condition specified in the query
\$addToSet	• Used to add array elements only if they are not present already
\$pop	• Used when the first or last item has to be removed from an array
\$pull	• Used when all the array elements that match the specified query condition has to be removed
\$push	• Used when elements must be added to an array
\$pullAll	• Used when all the array values that match the value specified in the query has to be removed

## \$pop Operator

The syntax for this operator is:

```
{ $pop: { <field>: <-1 | 1>, ... } }
```

In this syntax, `-1` is used to remove the first element and `1` is used to remove the last element in the array.

For example, consider that in the `Studentmarks` collection, the user wants to remove the first element from the `marks` array where `name` is "Oliver". To do so, the user can execute a query that uses the `$pop` operator as:

```
db.Studentmarks.updateOne( { "name": "Oliver" }, { $pop: { marks: -1 } } )
```

Figure 3.28 shows the output of this query.

```
Student_detail> db.Studentmarks.updateOne( { "name": "Oliver" }, { $pop: { marks: -1 } } )
{
  acknowledged: true,
  insertedId: null,
  matchedCount: 1,
  modifiedCount: 1,
  upsertedCount: 0
}
```

**Figure 3.28: Output of Query with \$pop Operator**

To view the updated document, the user can execute the query as:

```
db.Studentmarks.find( { "name": "Oliver" } )
```

Figure 3.29 shows the output of this query.

```
Student_details> db.Studentmarks.find( { "name": "Oliver" } )
[
  {
    _id: ObjectId("646f5c0760c2b086a2c39e2a"),
    name: 'Oliver',
    age: 17,
    sub: 1,
    marks: [ 85, 89 ],
    credit: 0
  }
]
```

**Figure 3.29: Updated Document**

## \$pull Operator

The syntax for this operator is:

```
{ $pull: { <field1>: <value|condition>,
           <field2>: <value|condition>, ... } }
```

Consider that in the Studentmarks collection, the user wants to remove the specific element from the marks array where name is "Oliver". To do so, the user can execute a query that uses the \$pull operator as:

```
db.Studentmarks.updateOne({ name: "Oliver" },
{ $pull: { marks: { $in: [85] } } })
```

Figure 3.30 shows the output of this query.

```
Student_detail> db.Studentmarks.updateOne({ name: "Oliver" }, { $pull: { marks: { $in: [85] } } })
{
  acknowledged: true,
  insertedId: null,
  matchedCount: 1,
  modifiedCount: 1,
  upsertedCount: 0
}
```

**Figure 3.30: Output of Query with \$pull Operator**

To view the updated document, the user can execute the query as:

```
db.Studentmarks.find({ "name": "Oliver" })
```

Figure 3.31 shows the output of this query.

```
Student_details> db.Studentmarks.find({ "name": "Oliver" })
[
  {
    _id: ObjectId("646f5c0760c2b086a2c39e2a"),
    name: 'Oliver',
    age: 17,
    sub: 1,
    marks: [ 89 ],
    credit: 0
  }
]
```

**Figure 3.31: Updated Document**

## \$push Operator

The syntax for this operator is:

```
{ $push: { <field1>: <value1>, ... } }
```

For example, consider that in the `Studentmarks` collection, the user wants to add a value to the `marks` array in the document where `name` is "Oliver". To do so, the user can execute a query that uses the `$push` operator as:

```
db.Studentmarks.updateOne( { "name": "Oliver" }, { $push: { marks: 99 } } )
```

Figure 3.32 shows the output of this query.

```
Student_detail> db.Studentmarks.updateOne( { "name": "Oliver" }, { $push: { marks: 99 } } )
{
  acknowledged: true,
  insertedId: null,
  matchedCount: 1,
  modifiedCount: 1,
  upsertedCount: 0
}
```

**Figure 3.32: Output of Query with \$push Operator**

To view the updated document, the user can execute the query as:

```
db.Studentmarks.find({ "name": "Oliver" })
```

Figure 3.33 shows the output of this query.

```
Student_details> db.Studentmarks.find({ "name": "Oliver" })
[
  {
    _id: ObjectId("646f5c0760c2b086a2c39e2a"),
    name: 'Oliver',
    age: 17,
    sub: 1,
    marks: [ 89, 99 ],
    credit: 0
  }
]
```

**Figure 3.33: Updated Document**

### **3.5 Other Operators**

---

Three other operators that are commonly used in MongoDB are:

**\$comment**

- Used to add comments to a query predicate to explain what the query is intended to do

**\$rand**

- Used to generate a random floating number between 0 and 1

**\$natural**

- Used to specify whether the documents must be iterated through in the natural order or the reverse order; takes value 1 for natural order and -1 for reverse order

#### **\$comment Operator**

The syntax for this operator is:

```
db.collection.find( { <query>, $comment: <comment> } )
```

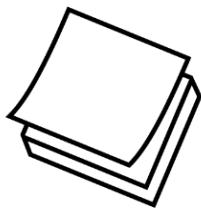
For example, consider that the user wants to add a comment for the document where the name is “Robert”. To do so, the user can execute a query that uses the \$comment operator as:

```
db.Studentmarks.find({ "name": "Robert",
$comment: "Displaying Robert mark details"})
```

Figure 3.34 shows the output of this query.

```
Student_details> db.Studentmarks.find({ "name": "Robert", $comment: "Displaying Robert mark details"})
[ {
  _id: ObjectId("646f5c0760c2b086a2c39e29"),
  name: 'Robert',
  age: 16,
  sub: 1,
  marks: [ 89, 78, 90 ],
  credit: 0
}]
```

**Figure 3.34: Output of the Query with \$comment Operator**



The comment is not visible in Figure 3.34. This is because the comments are stored in the `system.profile` collection. This collection is automatically created in the MongoDB server if database profiling is enabled. To view the comment, user can use the `find()` query on the `system.profile` collection.

## \$rand Operator

The syntax for this operator is:

```
{ $rand: {} }
```

For example, consider that in the `Studentmarks` collection, the user updates the field named `unit`. The value in this field will be automatically generated by the `$rand` operator. The value generated should be multiplied by 100 and rounded off to the highest integer. The query to achieve this objective is:

```
db.Studentmarks.updateMany( {}, [ { $set: { unit: { $ceil: { $multiply: [{ $rand: {} }, 100] } } } } ] )
```

The empty update filter matches every document in the collection. Figure 3.35 shows the output of this query:

```
Student_detail> db.Studentmarks.updateMany( {}, [ { $set: { unit: { $ceil: { $multiply: [{ $rand: {} }, 100] } } } } ] )  
{ acknowledged: true,  
  insertedId: null,  
  matchedCount: 4,  
  modifiedCount: 4,  
  upsertedCount: 0  
}
```

**Figure 3.35: Output of the Query with \$rand Operator**

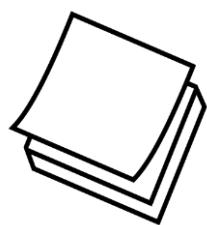
To view the updated documents, the user can execute the query:

```
db.Studentmarks.find()
```

Figure 3.36 shows the output of this query.

```
Student_details> db.Studentmarks.find()
[
  {
    _id: ObjectId("646f5c0760c2b086a2c39e29"),
    name: 'Robert',
    age: 16,
    sub: 1,
    marks: [ 89, 78, 90 ],
    credit: 0,
    unit: 55
  },
  {
    _id: ObjectId("646f5c0760c2b086a2c39e2a"),
    name: 'Oliver',
    age: 17,
    sub: 1,
    marks: [ 89, 99 ],
    credit: 0,
    unit: 50
  },
  {
    _id: ObjectId("646f5c0760c2b086a2c39e2b"),
    name: 'Henry',
    age: 17,
    sub: 1,
    marks: [ 88, 89, 76 ],
    credit: 0,
    unit: 21
  },
  {
    _id: ObjectId("646f5c0760c2b086a2c39e2c"),
    name: 'David',
    age: 16,
    sub: 2,
    marks: [ 86, 84, 66 ],
    Exam_Date: ISODate("2023-05-25T13:02:05.079Z"),
    unit: 37
  }
]
```

**Figure 3.36: Output of the `find()` Query**



---

The `unit` field is updated with the random number generated for all the documents in the collection. Note that the `unit` field in the document named "David" is updated with the random number as the `credit` field was renamed to `unit` using the `rename operator`. For all the other records, a new field named `unit` is added and updated with the random number.

---

### **3.6 Summary**

---

- MongoDB offers multiple operators to allow users to create complex queries.
- Main types of operators in MongoDB are query, projection, and update operators.
- Query operators allow operations that are based on comparison (both field and array), logic, and existence of a field.
- Projection operators help to specify what fields of a document a query must return.
- Update operators allow modification of both the field and array data.

## **Test Your Knowledge**

---

1. Which of the following comparison operator will allow you to list only the documents that contain a field whose value matches with any value in the specified array?
  - a. \$in
  - b. \$nin
  - c. \$ne
  - d. \$ni
  
2. Consider that `marks` is a collection which consists of a field `score`. Which of the following option will you use to find only the score value of datatype string?
  - a. `db.marks.find( { "Score" : { $type : 1 } })`
  - b. `db.marks.find( { "Score" : { $type : 2 } })`
  - c. `db.marks.find( { "Score" : { $type : 16 } })`
  - d. `db.marks.find( { "Score" : { $type : 3 } })`
  
3. Consider that `marks` is a collection that consists of a field `subject_marks` which is an array. Which of the following option will list documents with `subject_marks` field of size 3?
  - a. `db.subject_marks.find({marks:{$size:3}})`
  - b. `db.marks.find({subject_marks:{size:3}})`
  - c. `db.marks.find({"$subject_marks":{$size:3}})`
  - d. `db.marks.find({subject_marks:{$size:3}})`
  
4. Which of the following projection operator will project the first element in an array that matches the query condition?
  - a. \$slice
  - b. \$first
  - c. \$
  - d. \$elementfirst
  
5. Which of the following syntax allows you to remove the first element in the `marks` array?
  - a. { \$pop: { marks: -1 } }
  - b. { \$pop: { marks: 1 } }
  - c. { \$pop: { marks: 0 } }
  - d. { \$pop: { marks: \$first } }

## **Answers to Test Your Knowledge**

---

1	a
2	b
3	d
4	c
5	a

## **Try it Yourself**

---

1. Create a database named Student and a collection named Stud\_mark.
2. Insert the following four documents into the stud\_mark collection.

```
[  
  { name: "Adam",  
    gender:"M",  
    subjects:["Java", "C", "Python"],  
    marks:[89,78,90],  
    average:85.6  
,  
  { name: "Franklin",  
    gender:"M",  
    subjects:["C", "VB", "Python"],  
    marks:[78,85,89],  
    average:84  
,  
  { name: "Michael",  
    gender:"M",  
    subjects:["Java", "PHP"],  
    marks:[88,89],  
    average:88.5  
,  
  { name: "Amelia",  
    gender:"F",  
    subjects:["Ruby", "C++"],  
    marks:[86,87],  
    average: 86.5  
,  
]  
]
```

Using the collection stud\_mark perform the following tasks:

3. Find only the documents where the average value is equal to 84.
4. Find only the documents where the average value is greater than 85.
5. Display only the documents where the subjects array contains either Java or C++.
6. View only the documents where the average is greater than or equal to 87 and average is less than or equal to 90.
7. View all the documents where the subjects array has the value Java.
8. Display only the documents where the first element in the marks array is less than 80.

9. Display the details of the student named Adam where the marks array has only the first element and the second element.
10. Add a new date field Date\_of\_exam which shows the current date only for the student named "Amelia".
11. Increase the average value by 2 for the student named "Franklin".
12. Rename the field Date\_of\_exam as Examination\_date.